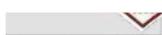


**A PROJECT REPORT**  
**On**  
**Laptop Price Prediction**

**Submitted by**

**Subham Shaw (10871022004)**  
**Nishant Yadav (10871022029)**  
**Angelina Deepshikha Hans (10871022009)**  
**Puja Verma (10871022011)**  
**Gaurav Kumar (10871022016)**

**Under the Guidance of**  
**Mr. Arnab Chakraborty**  
**Professor**



**Computer Applications**

**Asansol Engineering College**

**Asansol**

**Affiliated to**  
**MAULANA ABUL KALAM AZAD UNIVERSITY OF TECHNOLOGY**

## Contents

<b>Sl. No.</b>	<b>Topic</b>	<b>Page No.</b>
1.	Acknowledgement	1
2.	Project Objective	2
3.	Project Scope	3
4.	Data Description	4-5
5.	Data Pre-Processing	6-7
6.	EDA	8-11
7.	Feature Engineering	12-18
8.	Model Building	19-25
9.	Test Dataset	26
10.	Screenshots	27-28
11.	Code	29-77

12. Future Scope of Improvements 78

13. Certificates 7-84

## Acknowledgement

I take this opportunity to express my profound gratitude and deep regards to my faculty, **Prof. Arnab Chakraborty** for his exemplary guidance, monitoring and constant encouragement throughout the course of this project. The blessing, help and guidance given by him time to time shall carry me a long way in the journey of life on which I am about to embark.

I am obliged to my project team members for the valuable information provided by them in their respective fields. I am grateful for their cooperation during the period of my assignment.

Gaurav Kumar

Shubham Shaw

Angelina Deepshikha Hans

Puja Verma

Nishant Yadav

## Project Objective

In this project we have a shortened ‘Laptop Price’ Dataset from Kaggle. In this dataset, the target attribute is the Price . So, in this project we need to do Regression based on the attributes present in our dataset and predict the price of laptops according to its specifications.

Our methodology for solving the problems in the given project is described below:

- Load the required dataset.
- Study the dataset.
- Describe the dataset.
- Visualise the dataset.
- Find out if the dataset needs to be pre-processed.
  - It will be determined on the basis of whether the dataset has null values or outliers or any such discrepancy that might affect the output of the models to be trained.
- If the dataset is required to be pre-processed, take the necessary steps to pre-process the data.
- Find out the principal attributes for training.
- Split the given dataset for training and testing purpose.
- Fit the previously split train data in the aforementioned 4 models.
- Calculate the accuracy of the 4 models and find out the classification reports.
- Plot the necessary graphs.
- Use each trained model to predict the outcomes of the given test dataset.
- Choose the best model among the 4 trained models bases on the R2 Score and Mean Absolute Error.

## Project Scope

The Scope of the project is to develop a machine learning model that can predict the price of a laptop based on its specifications. This model will be helpful for various purposes such as:

- **Consumers:** get a better understanding of the fair market value of a laptop they are considering purchasing.
- **Retailers:** optimize their pricing strategies for laptops in their inventory.
- **Manufacturers:** gain insights into factors that influence laptop prices and make informed decisions about product development and pricing.

## Data Description

**Source of the data:** Kaggle. The given dataset is a shortened version of the original dataset in Kaggle.

**Data Description:** The given train dataset has 1303 rows and 12 columns

Attributes	Description	Data Type
Company	The manufacturer of the laptop (e.g., Apple, HP, Asus)	Categorical
TypeName	The type of laptop (e.g., Ultrabook, Notebook)	Categorical
Inches	The size of the laptop screen in inches (e.g., 13.3, 15.6)	Numeric
Screen Resolution	The resolution of the laptop screen (e.g., 1920x1080, 2560x1600)	Categorical
CPU	The central processing unit (CPU) of the laptop (e.g., Intel Core i5, AMD A9-Series)	Categorical
RAM	The amount of random access memory (RAM) in the laptop (e.g., 4GB, 8GB, 16GB)	Numeric
Memory	The type and storage capacity of the laptop's memory (e.g., 128GB SSD, 500GB HDD)	Categorical
GPU	The graphics processing unit (GPU) of the laptop (e.g., Intel Iris Plus Graphics, Nvidia GeForce MX150)	Categorical
OpSys	The operating system installed on the laptop (e.g., macOS, Windows 10)	Categorical
Weight	The weight of the laptop in kilograms (kg)	Numeric
Price	The price of the laptop (presumably in INR)	Numeric

Table 1: Data Description

Now we will pre-process the data. The methodology followed is given below:

- Checking for null values.
  - If null values are present, we will fill them or drop the row containing the null value based on the dataset.
- Checking for outliers.
  - If outliers are present, they will either be removed or replaced by following a suitable method depending on the dataset.

Null Values:

```
Unnamed: 0      0
Company         0
TypeName        0
Inches          0
ScreenResolution 0
Cpu             0
Ram             0
Memory          0
Gpu             0
OpSys           0
Weight          0
Price           0
dtype: int64
```

## Data Pre-processing

As the given dataset had Categorical and Non-categorical data mixed, we converted the categorical data into non-categorical data accordingly. We converted the String into Features. We converted the other categorical attributes into suitable numerical values. The following table shows the conversion record:

Non-Numeric to Numeric Change table		
<u>Column</u>	<u>Initial Value</u>	<u>Replaced Value</u>
Ram	8 GB	8
Weight	2 Kg	2
Touchscreen	NA	0
	NA	1
PPI	NA	Continuous Variable
Processor	NA	Categorical Values

Table 3: Feature Engineering

Now we have successfully handled Null values and converted non-numeric values to Numeric values. We didn't drop the rows with null values as we have a small dataset (only 1300 entries).

So, we are moving on to find if there are any outliers in our data and find the correlations of different attributes to our target i.e. 'Price' column in the dataset.

The following table gives the correlation value of each attribute with our target attribute i.e. ‘Price’:

Ram	0.742905
Weight	0.209867
Price	1.000000
Touchscreen	0.192917
Ips	0.253320
ppi	0.475368
HDD	-0.096891
SSD	0.670660
Name: Price, dtype: float64	

---

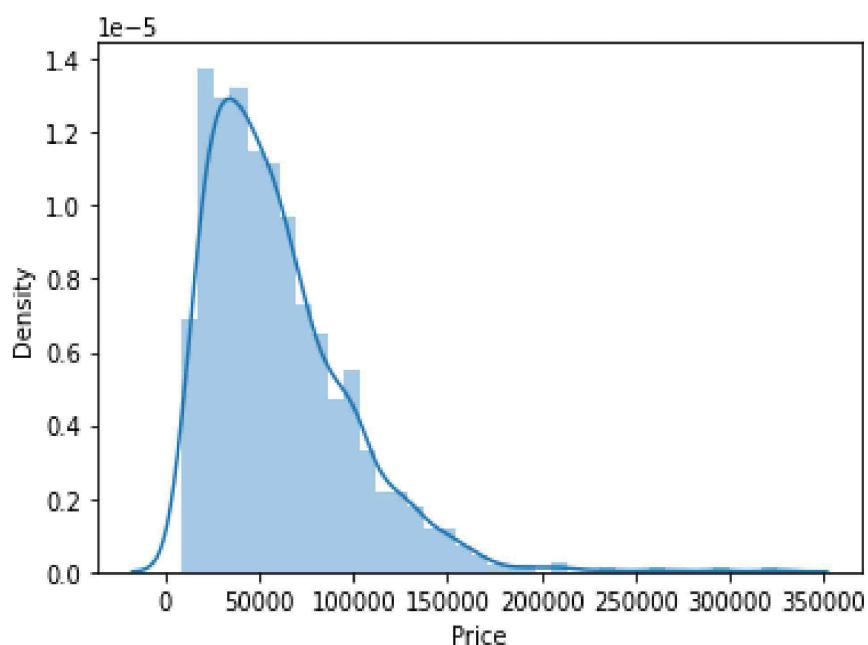
Table 5: Correlation values with Target Attribute

## EDA

Exploratory analysis is a process to explore and understand the data and data relationship in a complete depth so that it makes feature engineering and machine learning modeling steps smooth and streamlined for prediction. EDA involves Univariate, Bivariate, or Multivariate analysis. EDA helps to prove our assumptions true or false. In other words, it helps to perform hypothesis testing. We will start from the first column and explore each column and understand what impact it creates on the target column. At the required step, we will also perform preprocessing and feature engineering tasks. Our aim in performing in-depth EDA is to prepare and clean data for better machine learning modeling to achieve high performance and generalized models. So let's get started with analyzing and preparing the dataset for prediction.

### **1) Distribution of target column**

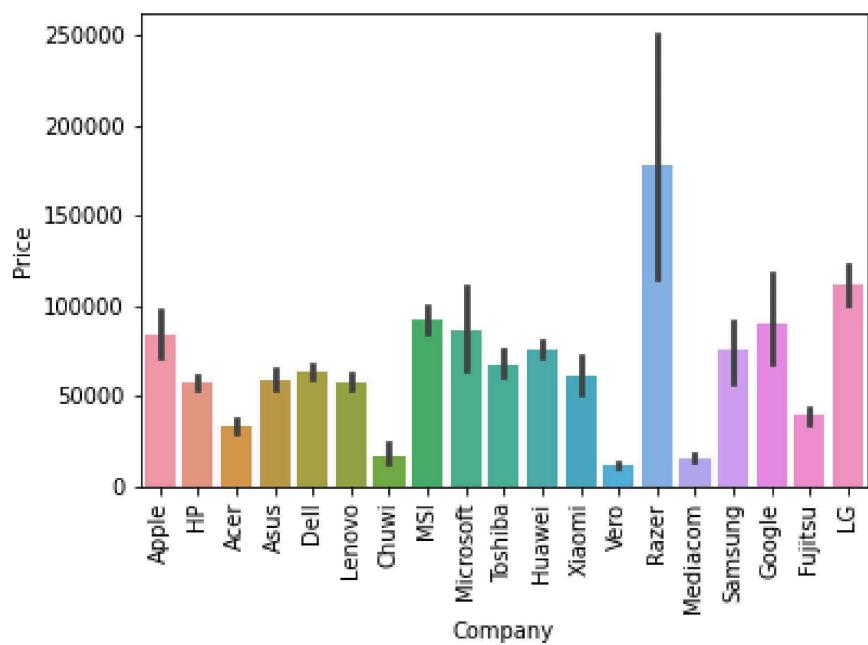
Working with regression problem statement target column distribution is important to understand.



## 2) Company column

we want to understand how does brand name impacts the laptop price or what is the average price of each laptop brand? If you plot a count plot(frequency plot) of a company then the major categories present are Lenovo, Dell, HP, Asus, etc.

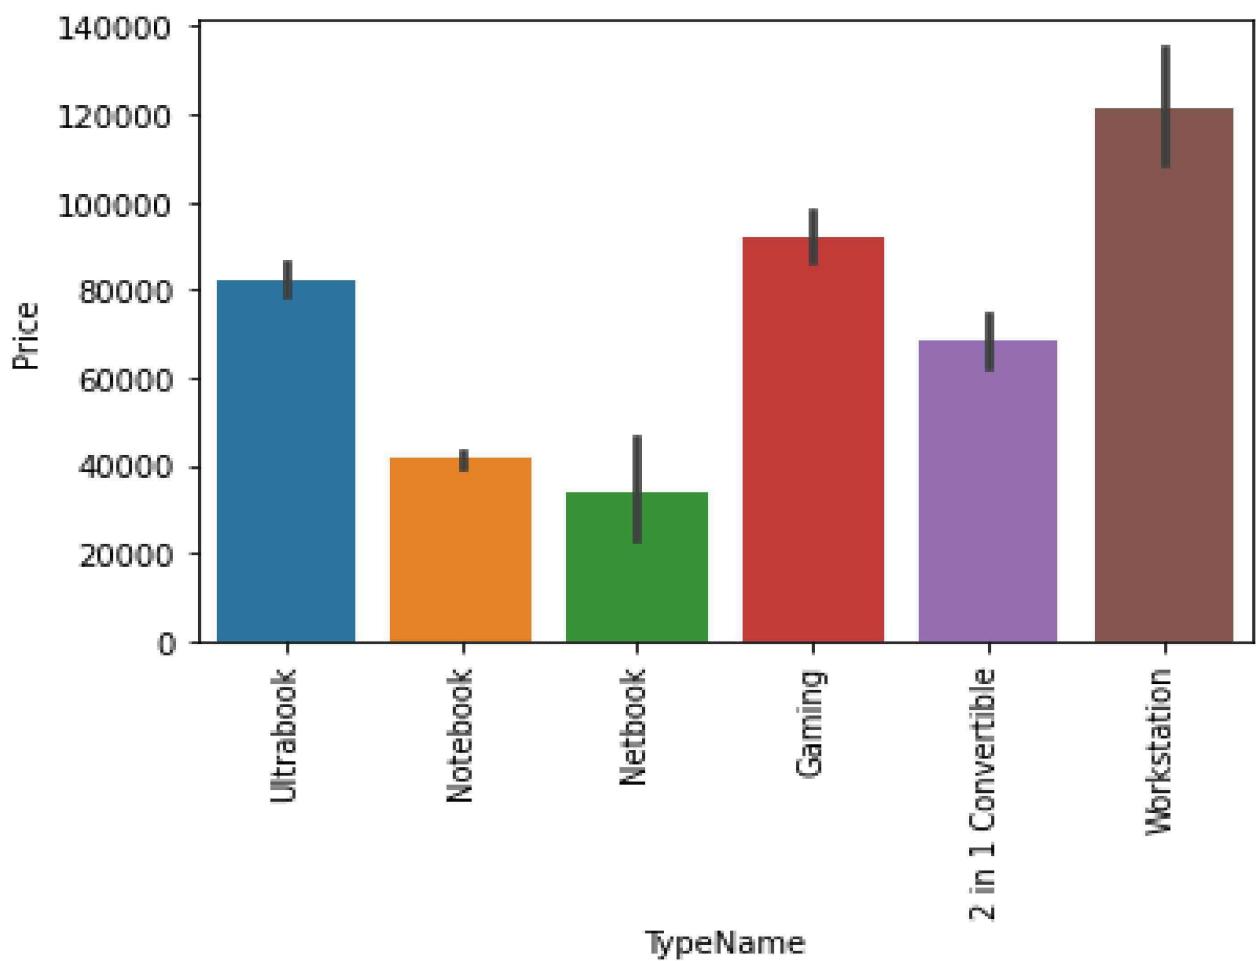
Now if we plot the company relationship with price then you can observe that how price varies with different brands.



Razer, Apple, LG, Microsoft, Google, MSI laptops are expensive, and others are in the budget range

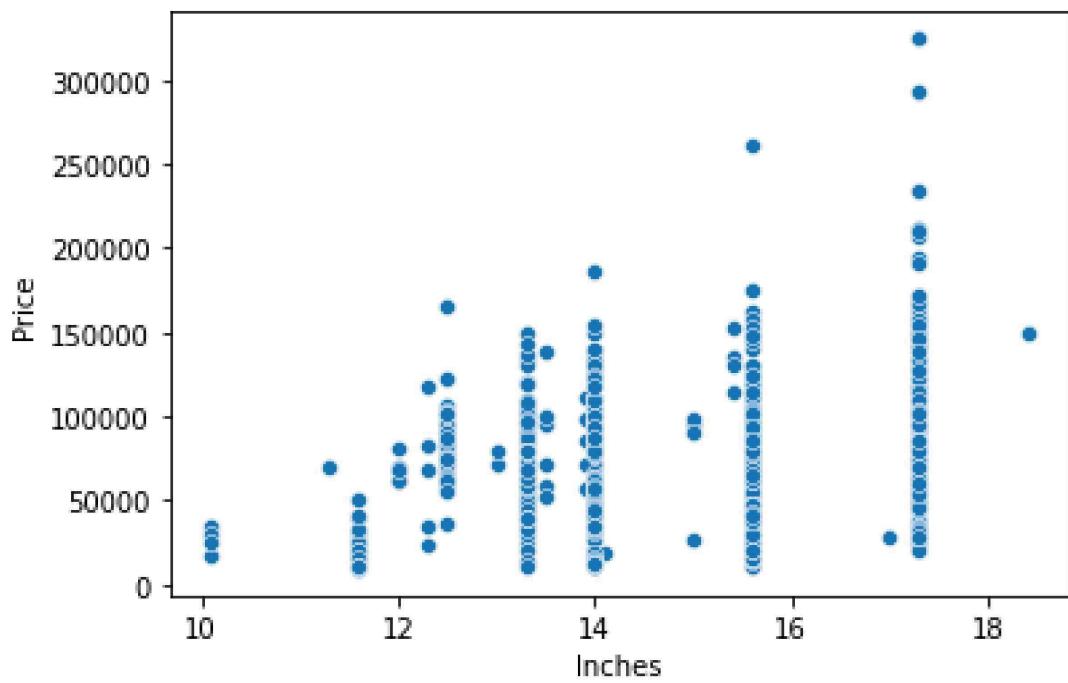
### 3) Type of laptop

Which type of laptop you are looking for like a gaming laptop, workstation, or notebook. As major people prefer notebook because it is under budget range and the same can be concluded from our data.



#### **4) Does the price vary with laptop size in inches?**

A Scatter plot is used when both the columns are numerical and it answers our question in a better way. From the below plot we can conclude that there is a relationship but not a strong relationship between the price and size column.



# Feature Engineering

Feature engineering is a process to convert raw data to meaningful information. there are many methods that come under feature engineering like transformation, categorical encoding, etc. Now the columns we have are noisy so we need to perform some feature engineering steps.

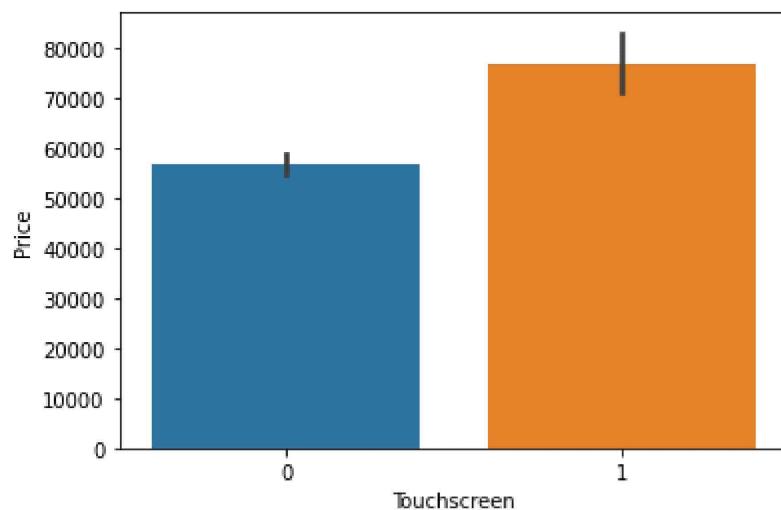
## ***Screen Resolution***

screen resolution contains lots of information. before any analysis first, we need to perform feature engineering over it. If you observe unique values of the column then we can see that all value gives information related to the presence of an IPS panel, are a laptop touch screen or not, and the X-axis and Y-axis screen resolution. So, we will extract the column into 3 new columns in the dataset.

## **Extract Touch screen information**

It is a binary variable so we can encode it as 0 and 1. one means the laptop is a touch screen and zero indicates not a touch screen.

If we plot the touch screen column against price then laptops with touch screens are expensive which is true in real life.



## **Extract IPS Channel presence information**

It is a binary variable and the code is the same we used above. The laptops with IPS channel are present less in our data but by observing relationship against the price of IPS channel laptops are high.

## **Extract X-axis and Y-axis screen resolution dimensions**

Now both the dimension are present at end of a string and separated with a cross sign. So first we will split the string with space and access the last string from the list. then split the string with a cross sign and access the zero and first index for X and Y-axis dimensions.

## **Replacing inches, X and Y resolution to PPI**

If you find the correlation of columns with price using the `corr` method then we can see that inches do not have a strong correlation but X and Y-axis resolution have a very strong resolution so we can take advantage of it and convert these three columns to a single column that is known as Pixel per inches(PPI). In the end, our goal is to improve the performance by having fewer features.

```
data.corr()['Price'].sort_values(ascending=False)
```

```
Price          1.000000
Ram           0.743007
X_res         0.556529
Y_res         0.552809
ppi           0.473487
Ips           0.252208
Weight        0.210370
Touchscreen   0.191226
Inches         0.068197
Name: Price, dtype: float64
```

So now we can drop the extra columns which are not of use. At this point, we have started keeping the important columns in our dataset.

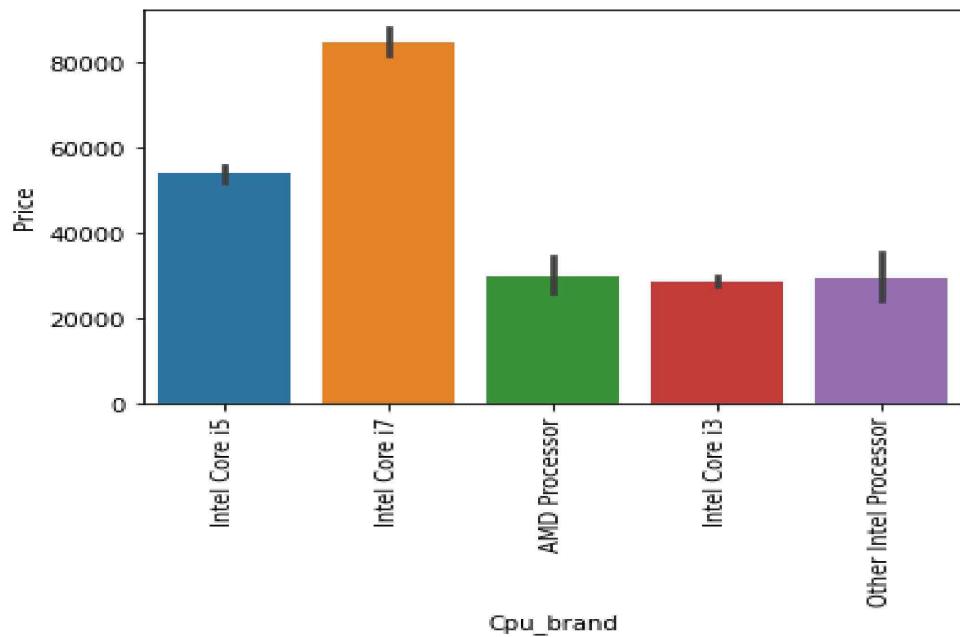
#### **CPU column**

If you observe the CPU column then it also contains lots of information. If you again use a unique function or value counts function on the CPU column then we have 118 different categories. The information it gives is about preprocessors in laptops and speed.

To extract the processor we need to extract the first three words from the string. we are having an Intel processor and AMD processor so we are keeping 5 categories in our dataset as i3, i5, i7, other intel processors, and AMD processors.

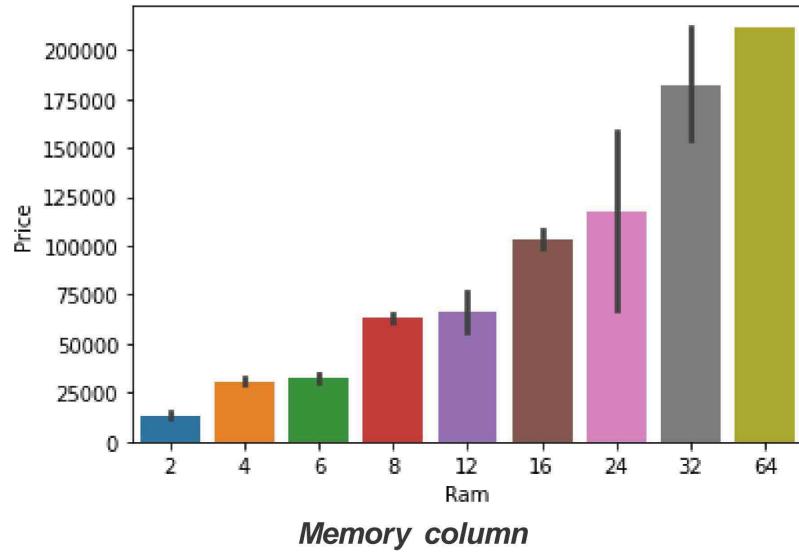
#### **How does the price vary with processors?**

we can again use our bar plot property to answer this question. And as obvious the price of i7 processor is high, then of i5 processor, i3 and AMD processor lies at the almost the same range. Hence price will depend on the processor



### **Price with Ram**

Again Bivariate analysis of price with Ram. If you observe the plot then Price is having a very strong positive correlation with Ram or you can say a linear relationship.



memory column is again a noisy column that gives an understanding of hard drives. many laptops came with HHD and SSD both, as well in some there is an external slot present to insert after purchase. This column can disturb your analysis if not feature engineer it properly. So If you use value counts on a column then we are having 4 different categories of memory as HHD, SSD, Flash storage, and hybrid.

First, we have cleaned the memory column and then made 4 new columns which are a binary column where each column contains 1 and 0 indicate that amount four is present and which is not present. Any laptop has a single type of memory or a combination of two. so in the first column, it consists of the first memory size and if the second slot is present in the laptop then the second column contains it else we fill the null values with zero. After that in a particular column, we have multiplied the values by their binary value. It means that if in any laptop particular memory is present then it contains binary value as one and the first value will be multiplied by it, and same with the second combination. For the laptop which does have a second slot, the value will be zero multiplied by zero is zero.

Now when we see the correlation of price then Hybrid and flash storage have very less or no correlation with a price. We will drop this column with CPU and memory which is no longer required.

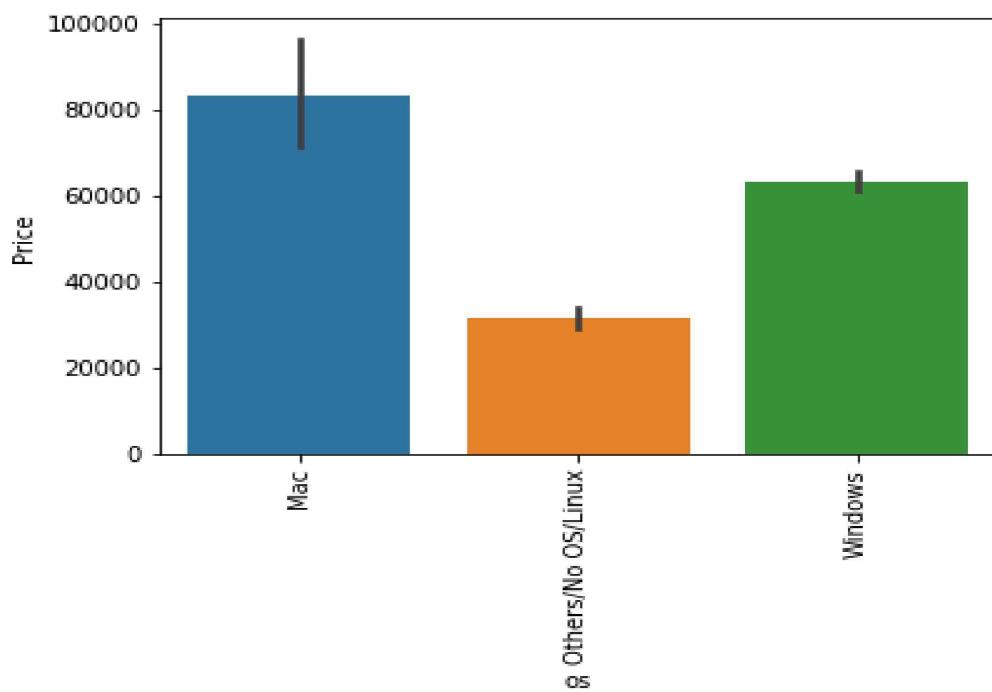
## **GPU Variable**

GPU(Graphical Processing Unit) has many categories in data. We are having which brand graphic card is there on a laptop. we are not having how many capacities like (6Gb, 12 Gb) graphic card is present. so we will simply extract the name of the brand.

## **Operating System Column**

There are many categories of operating systems. we will keep all windows categories in one, Mac in one, and remaining in others. This is a simple and most used feature engineering method, you can try something else if you find more correlation with price.

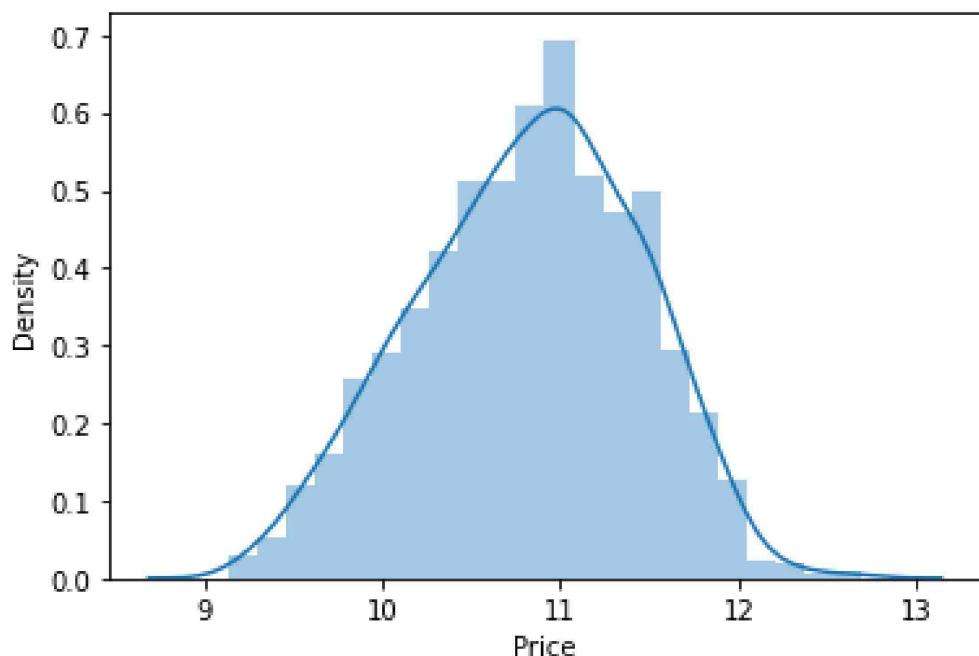
when you plot price against operating system then as usual Mac is most expensive.



## **Log-Normal Transformation**

we saw the distribution of the target variable above which was right-skewed. By transforming it to normal distribution performance of the algorithm will increase. we take the log of values that transform to the normal

distribution which you can observe below. So while separating dependent and independent variables we will take a log of price, and in displaying the result perform exponent of it.



## Model Building

Splitting data for training and testing purpose

We split the given train dataset into two parts for training and testing purpose. The split ratio we used is 0.75 which indicates we used 75% data for training purpose and 25% data for testing purpose. We will be using the same split ratio for all the models trained.

Random Forest Regressor Model

The object description of the Random Forest Classifier used is given below:

Object Name	RandomForestRegressor
Parameters	Value
max_depth	15
max_features	0.75
min_sample	0.5
n_estimators	100
random_state	100

Table 6: Object Parameter Table for Random Forest Regressor

Now we calculated R2 Score and MAE:

```
R2 score 0.8873402378382488  
MAE 0.15860130110457718
```

Now we will be training our required models. Our project goal requires us to train specific 4 classifier models viz.

1. KNN Regressor
2. Random Forest Regression
3. Decision Tree and
4. Regression (As the target attribute is Continous, We will use Linear Regression)

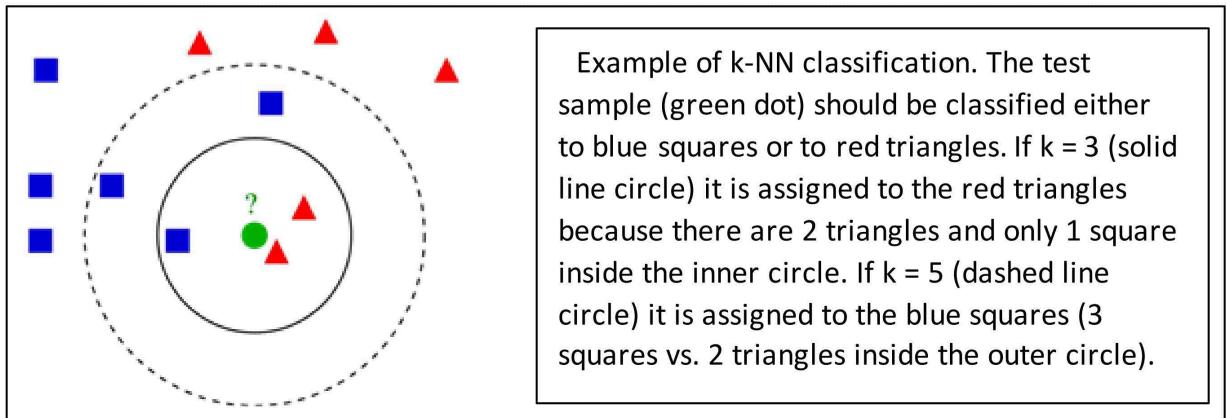
We will be using the final dataset obtained after pre-processing the given train dataset to train our required models.

#### KNN Regression

k-NN can be used for both classification and regression predictive problems. However, it is more widely used in classification problems in the industry. In pattern recognition, the k-nearest neighbors algorithm (k-NN) is a non-parametric method used for classification and regression. In both cases, the input consists of the k closest training examples in the feature space. The output depends on whether k-NN is used for classification or regression.

K-Nearest Neighbors (KNN) is a non-parametric machine learning algorithm that can be used for both classification and regression tasks. In the context of regression, KNN is often referred to as “K-Nearest Neighbors Regression” or “KNN Regression.” It’s a simple and intuitive algorithm that makes predictions by finding the K nearest data points to a given input and averaging their target values (for numerical regression) or selecting the majority class (for classification).

An example of k-NN classification:



The object description of the k-NN Classifier used is given below:

Object Name	KNeighborsClassifier
Parameters	Value
algorithm	'auto'
leaf_size	30
metric	'minkowski'
metric_params	None
n_neighbors	3
weights	'uniform'

Table 11: Object Parameter Table for k-NN Regressor

Now we calculate R2 Score and MAE for our trained k-NN model which is given below:

```
R2 score 0.803148868705085  
MAE 0.19264883332948868
```

Fig 7: R2 Score and MAE of k-NN Regressor

## Decision Tree

A Decision Tree Regressor is a supervised learning technique used for regression tasks. It works by creating a tree-like model that predicts a continuous value (in this case, price) based on a set of features (laptop specifications). Here's a breakdown of how it functions:

### 1. Building the Tree:

- The algorithm starts with the entire dataset and identifies the feature that best splits the data into subsets with the most significant difference in the target variable (price). This "best" feature is typically the one that minimizes a cost function, like variance in price.
- Each split creates a new decision node in the tree. The algorithm continues splitting the data at each node based on the best feature and its corresponding threshold value. This process repeats recursively until a stopping criteria is met, like reaching a maximum depth for the tree or having a sufficiently homogeneous subset of data points (similar prices) at a particular node.

### 2. Prediction:

- When presented with a new unseen laptop with its specifications, the model traverses the decision tree starting from the root node.
- At each node, the model compares the new data point's feature value with the threshold value associated with that node.
- Based on the comparison, the model follows a specific branch of the tree leading to a leaf node.
- The leaf node contains the predicted price for the new laptop. The prediction is typically the average price of the data points that reached that particular leaf node during training.

### Advantages of Decision Tree Regressor:

- Easy to interpret: The tree structure provides a clear visualization of how features influence the predicted price.
- Handles both categorical and numerical features: No need for separate data preprocessing for different feature

types.

- Robust to outliers: Less sensitive to outliers in the data compared to some other models.

#### **Disadvantages of Decision Tree Regressor:**

- Prone to overfitting: If the tree is allowed to grow too deep, it may become overly complex and capture noise in the training data, leading to poor performance on unseen data.
- High variance: Small changes in the training data can lead to significant variations in the tree structure and predictions.

The object description of the Decision Tree Regressor used is given below:

<b>Object Name</b>	<b>DecisionTreeRegressor</b>
<b>Parameters</b>	<b>Value</b>
max_depth	8

Table 17: Object Parameter Table for Decision Tree Regressor

Now, We Calculated the R2 Score and MAE of the Decision Tree Model:

```
R2 score 0.8411948861111318  
MAE 0.18163378208451647
```

Logistic Regression

Linear regression is a statistical method used to model the relationship between a dependent variable (also known as response variable or target variable) and one or more independent variables (also known as explanatory variables or predictor variables). It creates a linear equation that attempts to estimate the dependent variable based on the values of the independent variables.

Here's a breakdown of the key aspects of linear regression:

- **Equation:** The linear equation in linear regression typically takes the form of  $y = mx + b$ , where:
  - $y$  represents the dependent variable (e.g., laptop price in your project).
  - $x$  represents the independent variable (e.g., RAM size).
  - $m$  is the slope of the line, which indicates the change in  $y$  for a unit change in  $x$ .
  - $b$  is the  $y$ -intercept, which represents the predicted value of  $y$  when  $x$  is zero (not always meaningful in real-world applications).
- **Assumptions:** Linear regression makes certain assumptions about the data, including linearity (the relationship between variables can be modeled by a straight line), normality (errors are normally distributed), homoscedasticity (variance of errors is constant across all values of  $x$ ), and independence (errors are independent of each other).

The object description of the Logistic Regression Classifier used is given below:

Parameter	Description	Values
fit_intercept	(boolean, default=True) Whether to calculate the intercept ( $b$ ) for the linear equation.	TRUE
copy_X	(boolean, default=True) Controls whether to copy the input data ( $X$ ) before fitting the model. Copying prevents accidental modification of the original data.	TRUE
n_jobs	(int, default=None) The number of CPU cores to use for parallel computations during fitting.	None
positive	(boolean, default=False) When set to True, forces the coefficients to be positive. This is only applicable for dense arrays and might not be used in your code.	FALSE

Table 20: Object Parameter Table for Linear Regression

## Comparison of the Models trained

We trained 4 models using the 4 algorithms viz.

1. k-Nearest Neighbour
2. Gaussian Naive Bayes
3. Decision Tree and
4. Logistic Regression

The 4 models had different accuracy. The comparison of the accuracies of the models are given below:

<b>Model</b>	<b>R2 Score</b>
k-Nearest Neighbour	0.8031
Random Forest	0.8873
Decision Tree	0.8411
Linear Regression	0.8073

Table 24: Accuracy Comparison table

The following bar graph shows the accuracy comparison in graphical way:

### R2 Score of Different Models

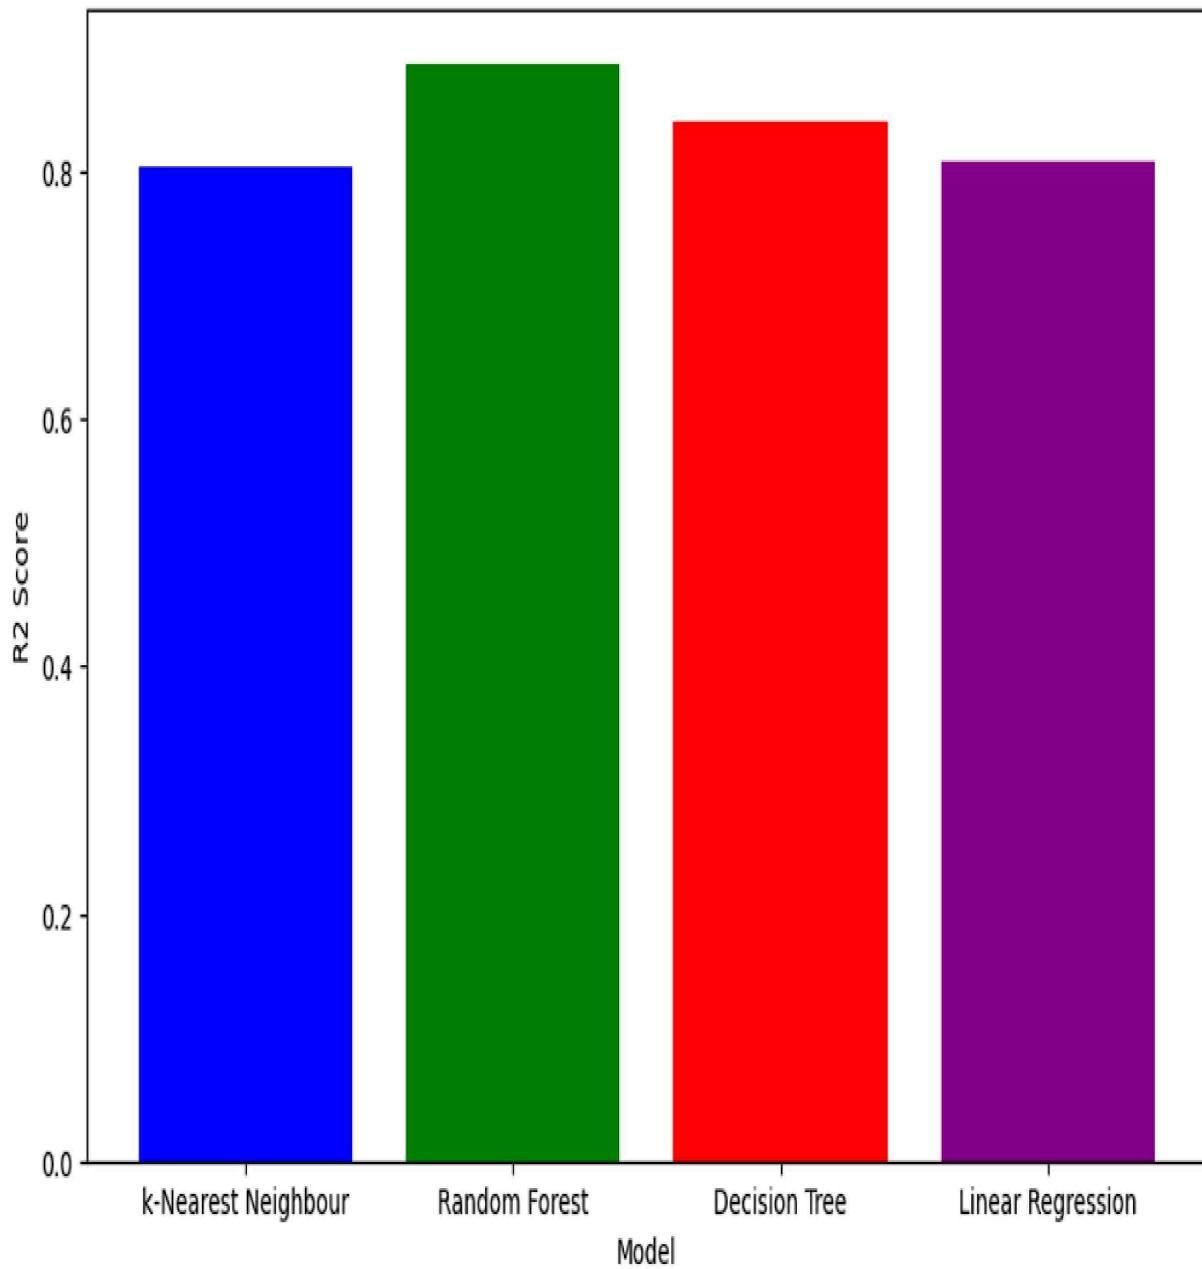


Fig 12: Comparison of accuracy of the 4 different models trained

Thus, from the above comparison we can see Random Forest Regressor has the highest R2 Score. So, our selected model is Random Forest

## Test Dataset

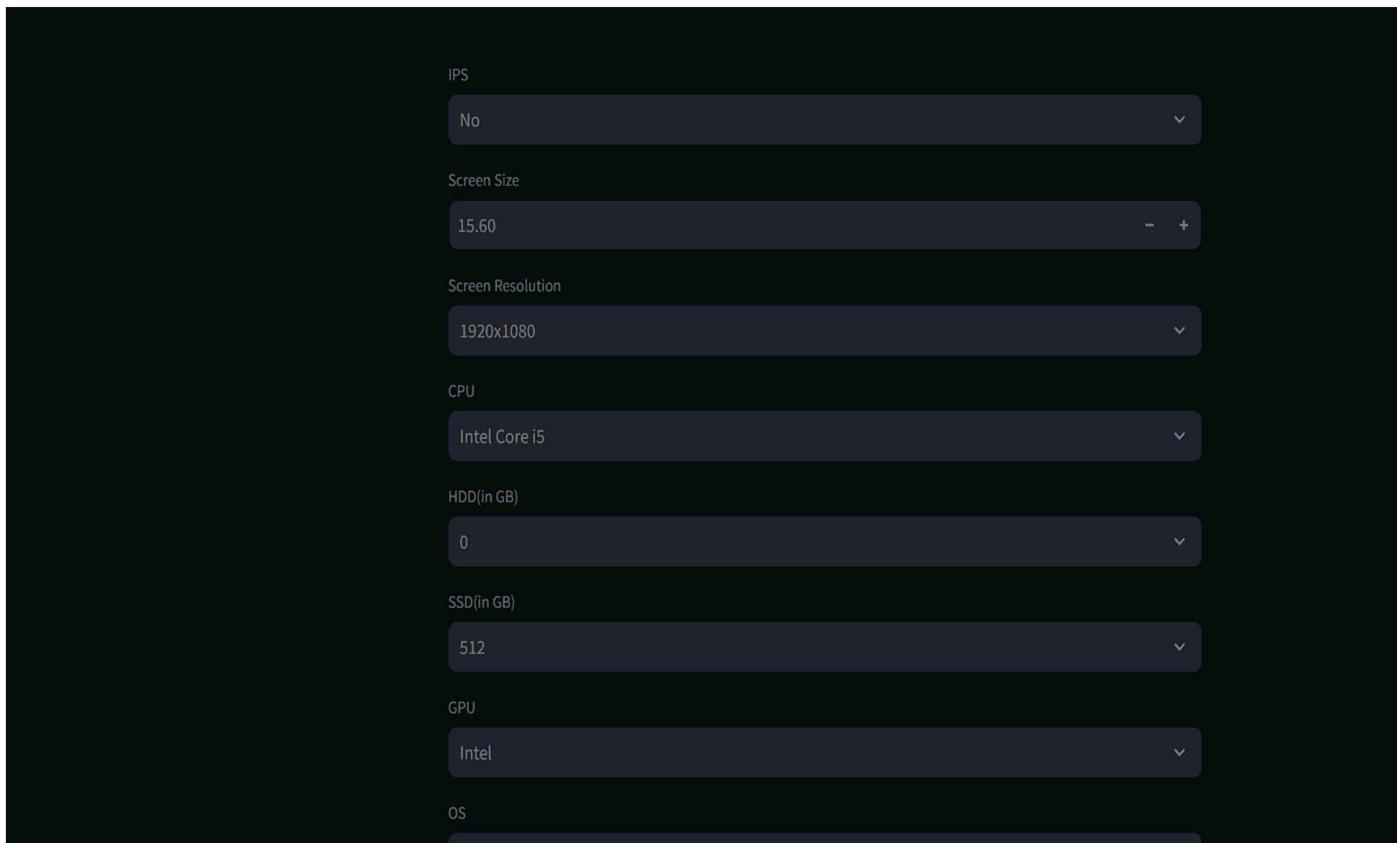
### Testing the given Test Dataset

We were given a test dataset for this laptop prediction problem. We pre-process the given test dataset in a similar way we pre-processed our train dataset. The methodology followed is given below:

- Change the categorical values to numeric values using the same process we used during changing values in our train dataset.
- Checking for null values.
  - If null values are present, we will fill them or drop the row containing the null value based on the dataset.
- Checking for outliers.
  - If outliers are present, they will either be removed or replaced by following a suitable method depending on the dataset.

After pre-processing of the test dataset had been done, we retrained the models with 100% of the train dataset given and used those trained models to predict the outcome for each input in the test dataset.

# SCREENSHOT



# Laptop Predictor

Brand

HP

Type

Notebook

RAM(in GB)

8

Weight of the Laptop

1.70

- +

Touchscreen

No

IPS

No

0

SSD(in GB)

512

GPU

Intel

OS

Windows

Predict Price

**The predicted price of this configuration is 69623**

## Code

### For GUI:

```
import streamlit as st  
import pickle  
import numpy as np  
  
# import the model  
pipe = pickle.load(open('pipe.pkl','rb'))  
df = pickle.load(open('df.pkl','rb'))  
  
st.title("Laptop Predictor")  
  
# brand  
company = st.selectbox('Brand',df['Company'].unique())  
  
# type of laptop  
type = st.selectbox('Type',df['TypeName'].unique())  
  
# Ram  
ram = st.selectbox('RAM(in GB)',[2,4,6,8,12,16,24,32,64])  
  
# weight  
weight = st.number_input('Weight of the Laptop')
```

```
# Touchscreen
touchscreen = st.selectbox('Touchscreen',['No','Yes'])

# IPS
ips = st.selectbox('IPS',['No','Yes'])

# screen size
screen_size = st.number_input('Screen Size')

# resolution
resolution = st.selectbox('Screen Resolution',['1920x1080','1366x768','1600x900','3840x2160','3200x1800','2880x1800','2560x1600','2560x1440','2304x1440'])

#cpu
cpu = st.selectbox('CPU',df['Cpu brand'].unique())

hdd = st.selectbox('HDD(in GB)',[0,128,256,512,1024,2048])

ssd = st.selectbox('SSD(in GB)',[0,8,128,256,512,1024])

gpu = st.selectbox('GPU',df['Gpu brand'].unique())

os = st.selectbox('OS',df['os'].unique())
```

```
if st.button('Predict Price'):

    # query

    ppi = None

    if touchscreen == 'Yes':
        touchscreen = 1
    else:
        touchscreen = 0

    if ips == 'Yes':
        ips = 1
    else:
        ips = 0

    X_res = int(resolution.split('x')[0])
    Y_res = int(resolution.split('x')[1])
    ppi = ((X_res*2) + (Y_res2))*0.5/screen_size
    query = np.array([company,type,ram,weight,touchscreen,ips,ppi
    ,cpu,hdd,ssd,gpu,os])

    query = query.reshape(1,12)

    st.title("The predicted price of this configuration is " +
    str(int(np.exp(pipe.predict(query)[0]))))
```

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

df = pd.read_csv('laptop_data.csv')
df.head()

      Unnamed: 0 Company    TypeName   Inches
ScreenResolution \
0            0     Apple  Ultrabook    13.3  IPS Panel Retina Display
2560x1600
1            1     Apple  Ultrabook    13.3
1440x900
2            2       HP  Notebook    15.6           Full HD
1920x1080
3            3     Apple  Ultrabook    15.4  IPS Panel Retina Display
2880x1800
4            4     Apple  Ultrabook    13.3  IPS Panel Retina Display
2560x1600

          Cpu    Ram           Memory \
0  Intel Core i5 2.3GHz  8GB  128GB SSD
1  Intel Core i5 1.8GHz  8GB  128GB Flash Storage
2  Intel Core i5 7200U 2.5GHz  8GB  256GB SSD
3  Intel Core i7 2.7GHz 16GB  512GB SSD
4  Intel Core i5 3.1GHz  8GB  256GB SSD

          Gpu  OpSys  Weight      Price
0  Intel Iris Plus Graphics 640  macOS  1.37kg  71378.6832
1  Intel HD Graphics 6000  macOS  1.34kg  47895.5232
2  Intel HD Graphics 620  No OS  1.86kg  30636.0000
3  AMD Radeon Pro 455  macOS  1.83kg  135195.3360
4  Intel Iris Plus Graphics 650  macOS  1.37kg  96095.8080

df.shape
(1303, 12)

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1303 entries, 0 to 1302
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        1303 non-null   int64  
 1   Company          1303 non-null   object  
 2   TypeName         1303 non-null   object  
 3   Inches           1303 non-null   float64
 4   ScreenResolution 1303 non-null   object 

```

```

5   Cpu           1303 non-null  object
6   Ram           1303 non-null  object
7   Memory        1303 non-null  object
8   Gpu           1303 non-null  object
9   OpSys         1303 non-null  object
10  Weight        1303 non-null  object
11  Price          1303 non-null  float64
dtypes: float64(2), int64(1), object(9)
memory usage: 76.4+ KB

df.duplicated().sum()

0

df.isnull().sum()

Unnamed: 0      0
Company         0
TypeName        0
Inches          0
ScreenResolution 0
Cpu             0
Ram             0
Memory          0
Gpu             0
OpSys           0
Weight          0
Price           0
dtype: int64

df.drop(columns=['Unnamed: 0'], inplace=True)

df.head()

   Company  TypeName  Inches           ScreenResolution \
0    Apple   Ultrabook    13.3  IPS Panel Retina Display  2560x1600
1    Apple   Ultrabook    13.3                               1440x900
2     HP    Notebook    15.6                  Full HD 1920x1080
3    Apple   Ultrabook    15.4  IPS Panel Retina Display  2880x1800
4    Apple   Ultrabook    13.3  IPS Panel Retina Display  2560x1600

                           Cpu   Ram       Memory \
0  Intel Core i5 2.3GHz  8GB  128GB SSD
1  Intel Core i5 1.8GHz  8GB  128GB Flash Storage
2  Intel Core i5 7200U 2.5GHz  8GB  256GB SSD
3  Intel Core i7 2.7GHz 16GB  512GB SSD
4  Intel Core i5 3.1GHz  8GB  256GB SSD

                           Gpu  OpSys  Weight      Price
0  Intel Iris Plus Graphics 640  macOS  1.37kg  71378.6832
1  Intel HD Graphics 6000  macOS  1.34kg  47895.5232

```

```

2           Intel HD Graphics 620  No OS  1.86kg  30636.0000
3           AMD Radeon Pro 455  macOS  1.83kg  135195.3360
4  Intel Iris Plus Graphics 650  macOS  1.37kg  96095.8080

df['Ram'] = df['Ram'].str.replace('GB', '')
df['Weight'] = df['Weight'].str.replace('kg', '')

df.head()

   Company    TypeName  Inches      ScreenResolution \
0   Apple     Ultrabook    13.3  2560x1600
1   Apple     Ultrabook    13.3  1440x900
2    HP       Notebook    15.6  Full HD 1920x1080
3   Apple     Ultrabook    15.4  2880x1800
4   Apple     Ultrabook    13.3  2560x1600

                           Cpu  Ram      Memory \
0     Intel Core i5 2.3GHz  8  128GB SSD
1     Intel Core i5 1.8GHz  8  128GB Flash Storage
2  Intel Core i5 7200U 2.5GHz  8  256GB SSD
3     Intel Core i7 2.7GHz  16  512GB SSD
4     Intel Core i5 3.1GHz  8  256GB SSD

          Gpu  OpSys  Weight      Price
0  Intel Iris Plus Graphics 640  macOS  1.37  71378.6832
1  Intel HD Graphics 6000  macOS  1.34  47895.5232
2  Intel HD Graphics 620  No OS  1.86  30636.0000
3  AMD Radeon Pro 455  macOS  1.83  135195.3360
4  Intel Iris Plus Graphics 650  macOS  1.37  96095.8080

df['Ram'] = df['Ram'].astype('int32')
df['Weight'] = df['Weight'].astype('float32')

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1303 entries, 0 to 1302
Data columns (total 11 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   Company           1303 non-null   object 
 1   TypeName          1303 non-null   object 
 2   Inches             1303 non-null   float64
 3   ScreenResolution  1303 non-null   object 
 4   Cpu                1303 non-null   object 
 5   Ram               1303 non-null   int32  
 6   Memory            1303 non-null   object 
 7   Gpu                1303 non-null   object 
 8   OpSys             1303 non-null   object 
 9   Weight            1303 non-null   float32
 10  Price             1303 non-null   float64

```

```
dtypes: float32(1), float64(2), int32(1), object(7)
memory usage: 66.2+ KB

%pip install -q seaborn
import seaborn as sns

sns.distplot(df['Price'])

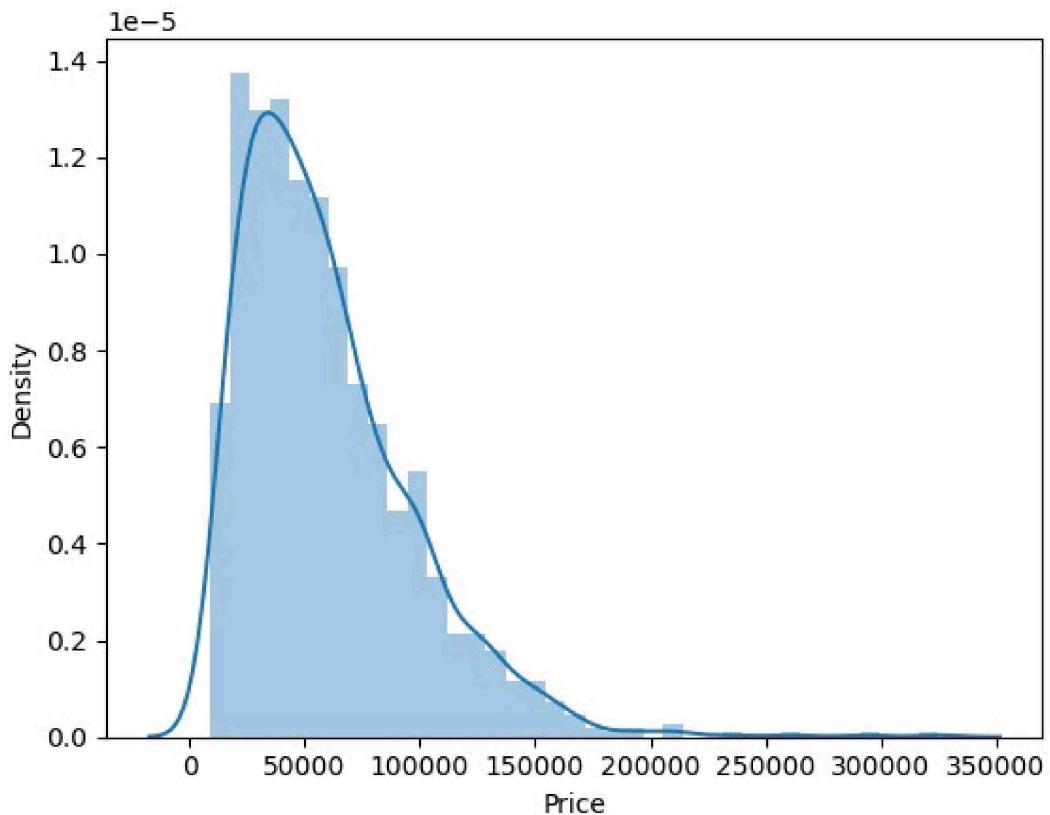
<ipython-input-18-87e11caeb2c4>:1: UserWarning:
`distplot` is a deprecated function and will be removed in seaborn
v0.14.0.

Please adapt your code to use either `displot` (a figure-level
function with
similar flexibility) or `histplot` (an axes-level function for
histograms).

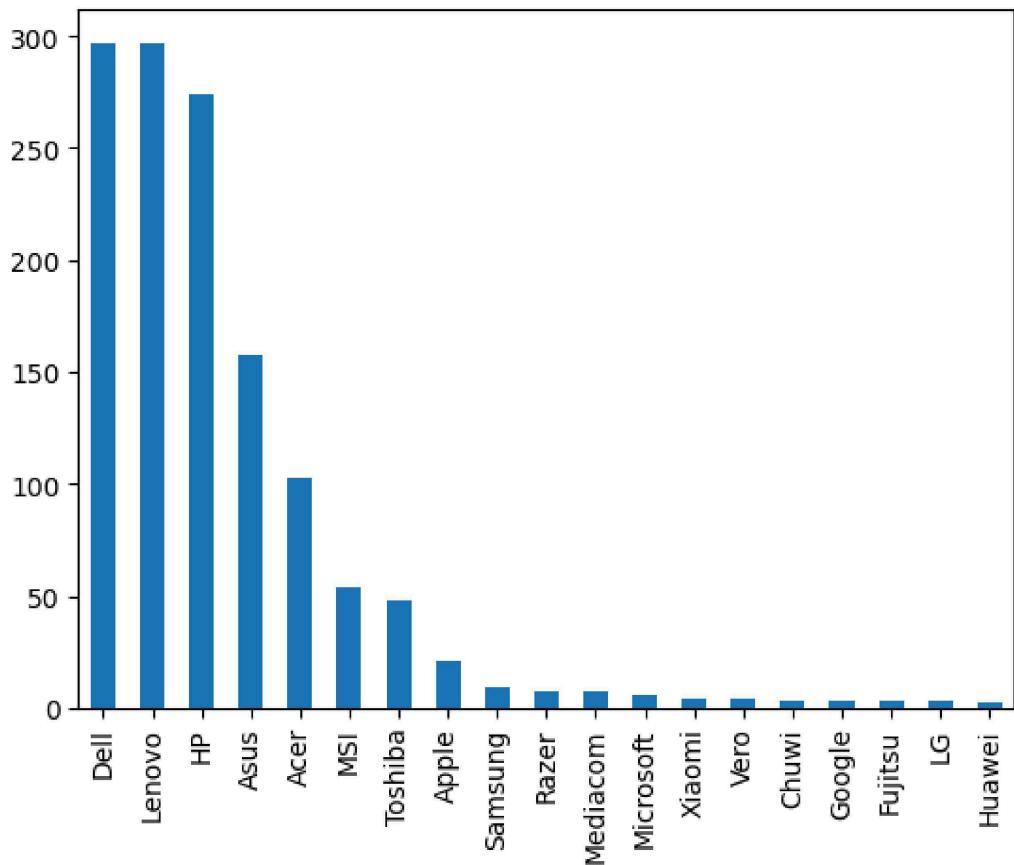
For a guide to updating your code to use the new functions, please see
https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751

sns.distplot(df['Price'])

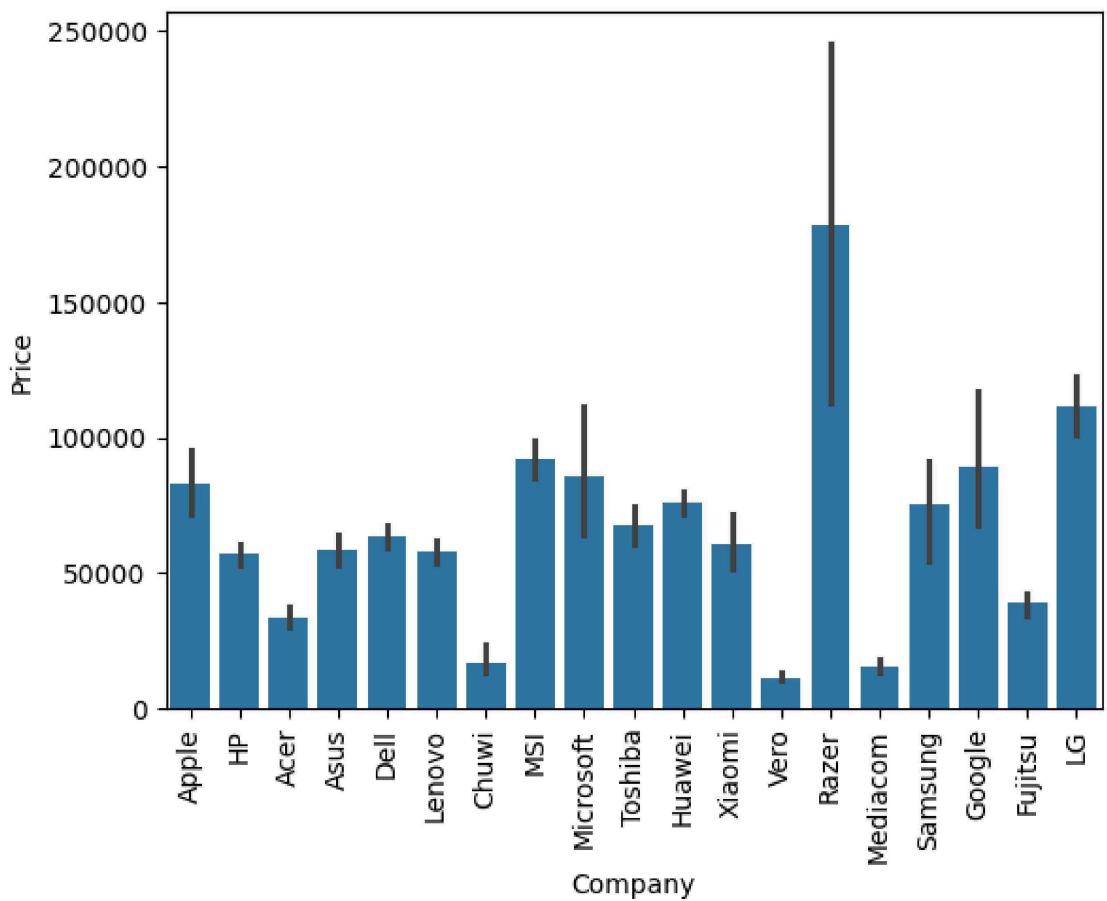
<AxesSubplot:xlabel='Price', ylabel='Density'>
```



```
df['Company'].value_counts().plot(kind='bar')  
<AxesSubplot:>
```

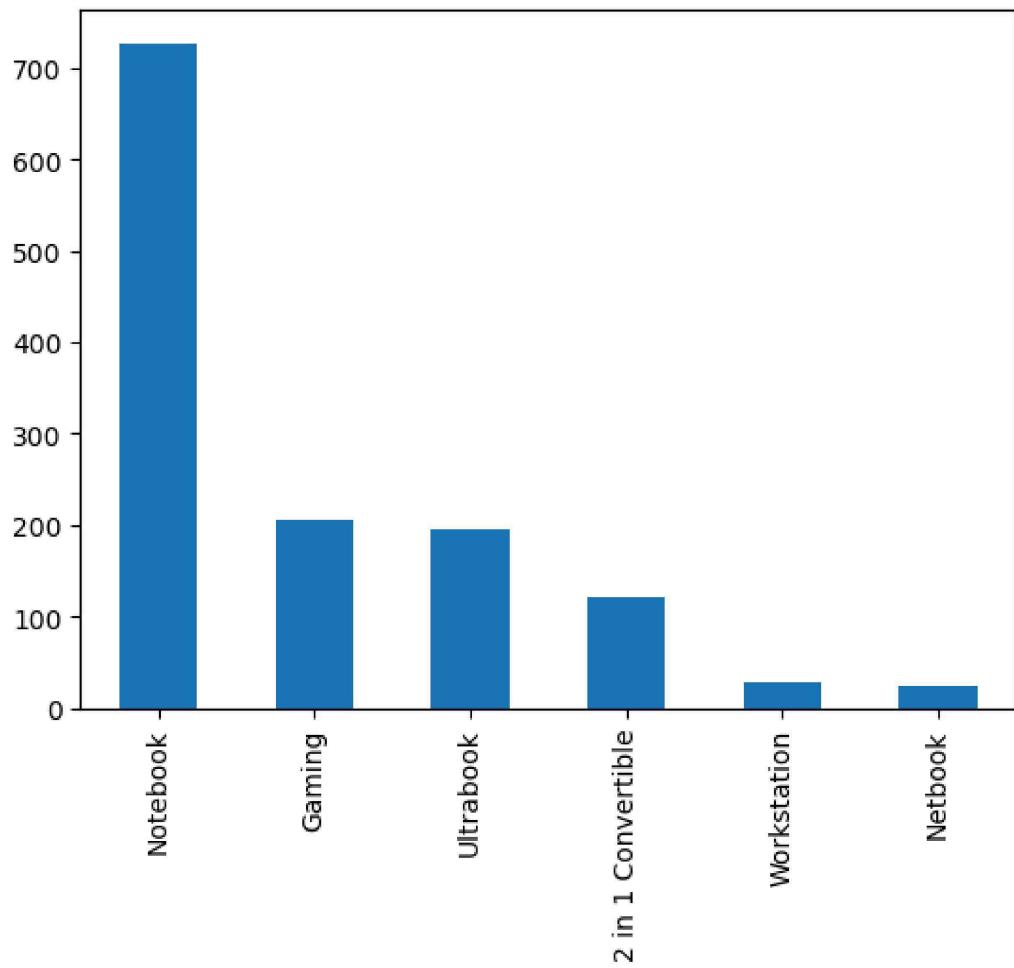


```
sns.barplot(x=df['Company'],y=df['Price'])  
plt.xticks(rotation='vertical')  
plt.show()
```

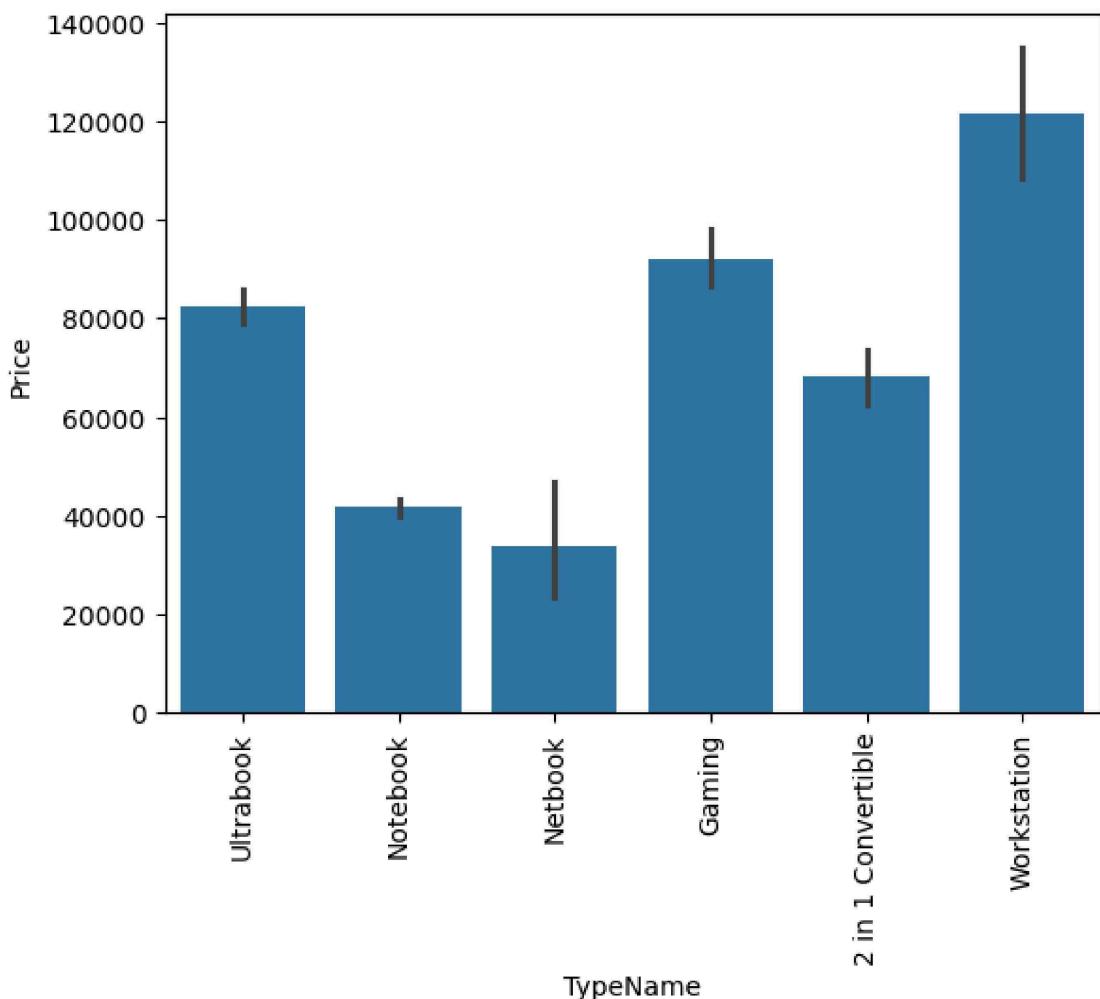


```
df[ 'TypeName' ].value_counts().plot(kind='bar')
```

```
<AxesSubplot:>
```



```
sns.barplot(x=df['TypeName'],y=df['Price'])
plt.xticks(rotation='vertical')
plt.show()
```



```
sns.distplot(df['Inches'])

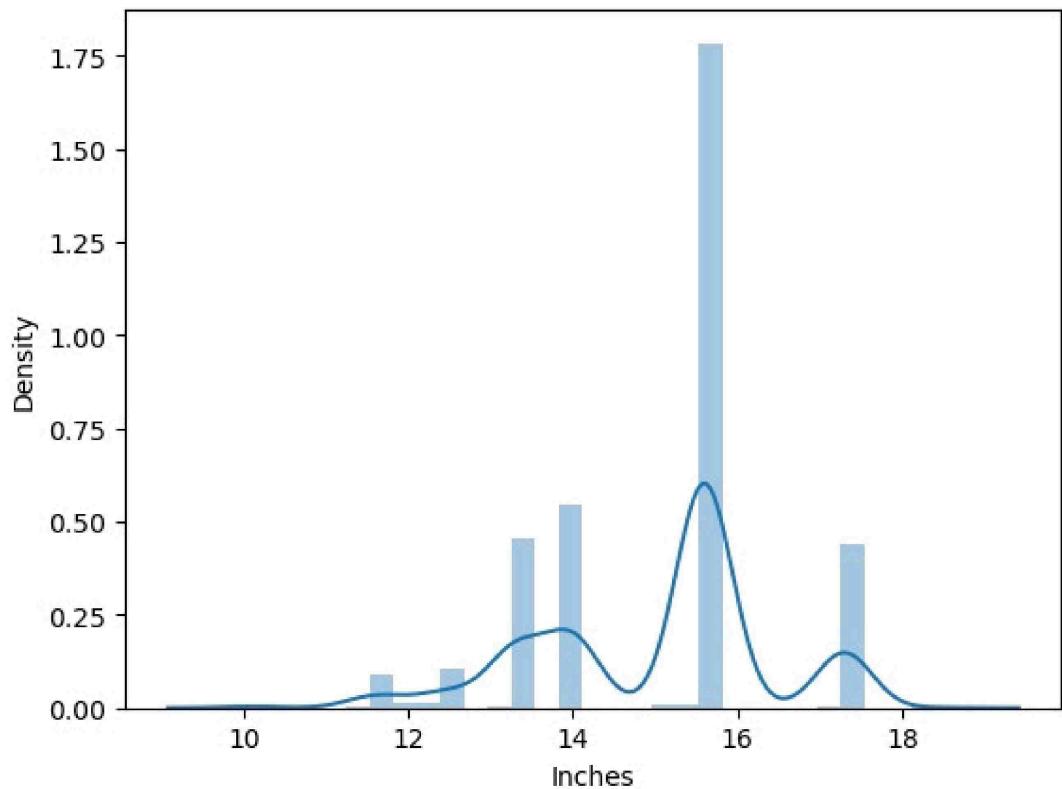
<ipython-input-23-51888cb550e6>:1: UserWarning:
`distplot` is a deprecated function and will be removed in seaborn
v0.14.0.

Please adapt your code to use either `displot` (a figure-level
function with
similar flexibility) or `histplot` (an axes-level function for
histograms).
```

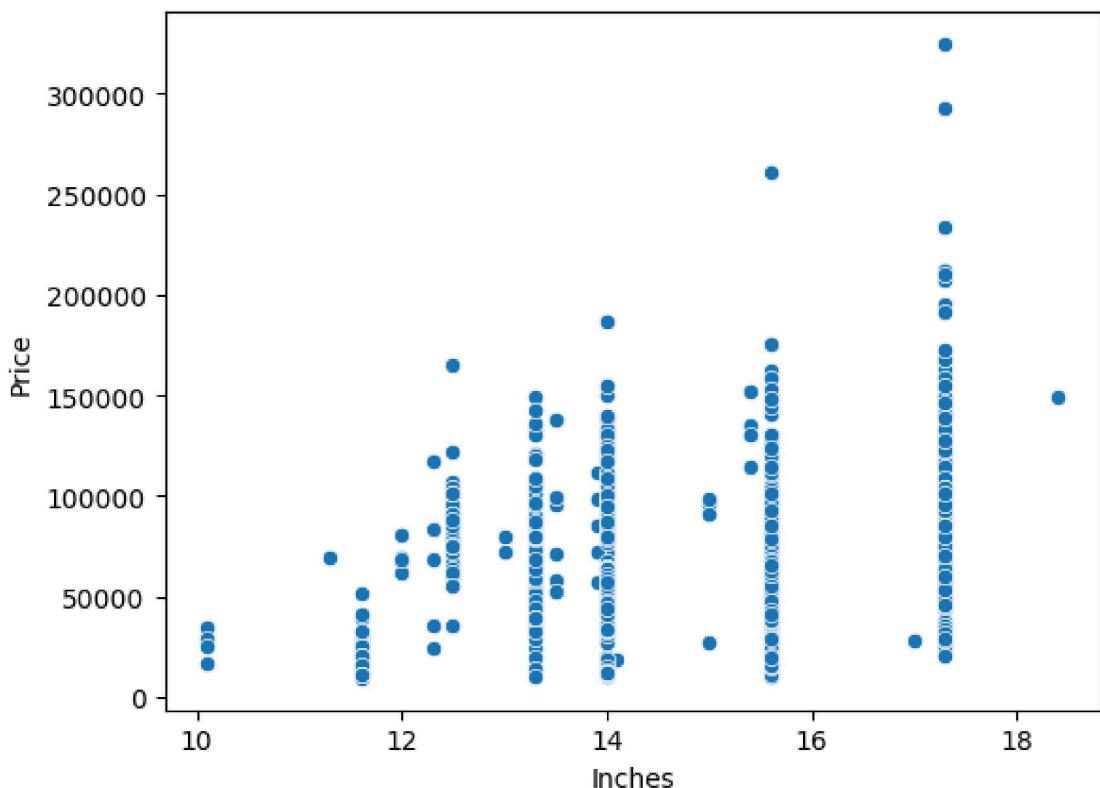
For a guide to updating your code to use the new functions, please see  
<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['Inches'])

<AxesSubplot:xlabel='Inches', ylabel='Density'>
```



```
sns.scatterplot(x=df['Inches'],y=df['Price'])  
<AxesSubplot:xlabel='Inches', ylabel='Price'>
```



```
df['ScreenResolution'].value_counts()
```

Full HD 1920x1080	507
1366x768	281
IPS Panel Full HD 1920x1080	230
IPS Panel Full HD / Touchscreen 1920x1080	53
Full HD / Touchscreen 1920x1080	47
1600x900	23
Touchscreen 1366x768	16
Quad HD+ / Touchscreen 3200x1800	15
IPS Panel 4K Ultra HD 3840x2160	12
IPS Panel 4K Ultra HD / Touchscreen 3840x2160	11
4K Ultra HD / Touchscreen 3840x2160	10
4K Ultra HD 3840x2160	7
Touchscreen 2560x1440	7
IPS Panel 1366x768	7
IPS Panel Quad HD+ / Touchscreen 3200x1800	6
IPS Panel Retina Display 2560x1600	6
IPS Panel Retina Display 2304x1440	6
Touchscreen 2256x1504	6
IPS Panel Touchscreen 2560x1440	5
IPS Panel Retina Display 2880x1800	4
IPS Panel Touchscreen 1920x1200	4
1440x900	4

```

IPS Panel 2560x1440                                4
IPS Panel Quad HD+ 2560x1440                      3
Quad HD+ 3200x1800                                3
1920x1080                                         3
Touchscreen 2400x1600                            3
2560x1440                                         3
IPS Panel Touchscreen 1366x768                     3
IPS Panel Touchscreen / 4K Ultra HD 3840x2160      2
IPS Panel Full HD 2160x1440                        2
IPS Panel Quad HD+ 3200x1800                      2
IPS Panel Retina Display 2736x1824                 1
IPS Panel Full HD 1920x1200                        1
IPS Panel Full HD 2560x1440                        1
IPS Panel Full HD 1366x768                         1
Touchscreen / Full HD 1920x1080                   1
Touchscreen / Quad HD+ 3200x1800                   1
Touchscreen / 4K Ultra HD 3840x2160                 1
IPS Panel Touchscreen 2400x1600                     1
Name: ScreenResolution, dtype: int64

df['Touchscreen'] = df['ScreenResolution'].apply(lambda x: 1 if
'Touchscreen' in x else 0)

df.sample(5)

      Company   TypeName   Inches          ScreenResolution \
553       HP    Notebook     17.3           1600x900
239       Acer   Gaming      15.6    IPS Panel Full HD 1920x1080
47        Asus   Gaming      17.3           Full HD 1920x1080
1034      HP    Notebook     15.6           Full HD 1920x1080
917       MSI   Gaming      17.3           Full HD 1920x1080

                           Cpu   Ram          Memory \
553   Intel Core i3 6006U 2GHz     8           1TB HDD
239   Intel Core i5 7300HQ 2.5GHz   8           256GB SSD
47        AMD Ryzen 1700 3GHz     8  256GB SSD + 1TB HDD
1034  Intel Core i7 6500U 2.5GHz   8           256GB SSD
917   Intel Core i7 7700HQ 2.8GHz  16  256GB SSD + 1TB HDD

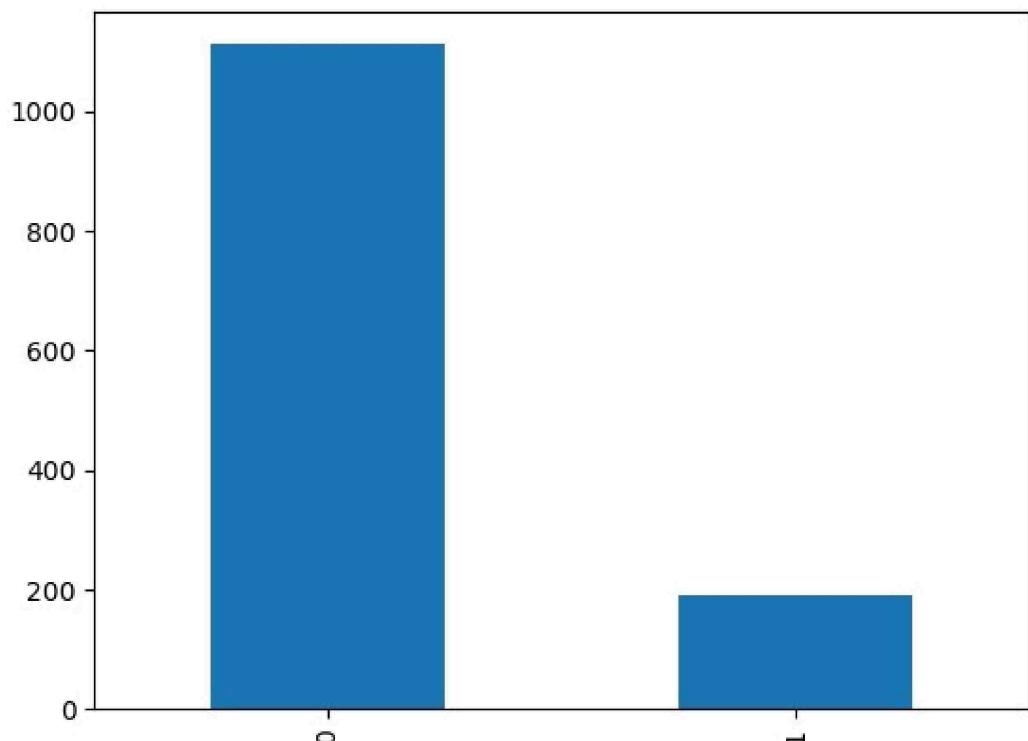
                    Gpu   OpSys  Weight   Price
Touchscreen
553       Intel HD Graphics 520 Windows 10    2.65  28992.3120
0
239       Nvidia GeForce GTX 1050 Windows 10    2.50  45074.8800
0
47        AMD Radeon RX 580 Windows 10    3.20  69210.7200
0
1034      Intel HD Graphics 520 Windows 10    1.84  65480.5872
0

```

```
917    Nvidia GeForce GTX 1050 Ti  Windows 10      2.70  85194.7200
0
```

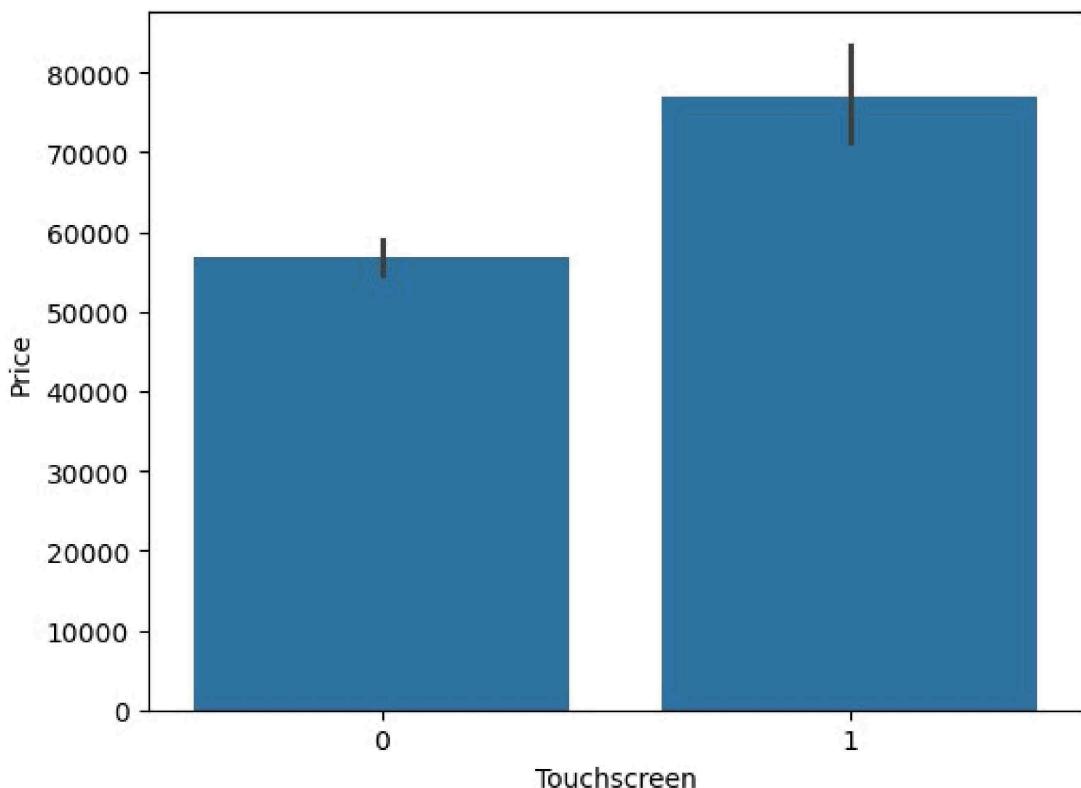
```
df['Touchscreen'].value_counts().plot(kind='bar')
```

```
<AxesSubplot:>
```



```
sns.barplot(x=df['Touchscreen'],y=df['Price'])
```

```
<AxesSubplot:xlabel='Touchscreen', ylabel='Price'>
```



```

df['Ips'] = df['ScreenResolution'].apply(lambda x:1 if 'IPS' in x else 0)

df.head()

    Company    TypeName  Inches      ScreenResolution \
0   Apple     Ultrabook  13.3  IPS Panel Retina Display 2560x1600
1   Apple     Ultrabook  13.3                               1440x900
2    HP       Notebook  15.6          Full HD 1920x1080
3   Apple     Ultrabook  15.4  IPS Panel Retina Display 2880x1800
4   Apple     Ultrabook  13.3  IPS Panel Retina Display 2560x1600

                           Cpu  Ram      Memory \
0        Intel Core i5 2.3GHz   8  128GB SSD
1        Intel Core i5 1.8GHz   8  128GB Flash Storage
2  Intel Core i5 7200U 2.5GHz   8  256GB SSD
3        Intel Core i7 2.7GHz  16  512GB SSD
4        Intel Core i5 3.1GHz   8  256GB SSD

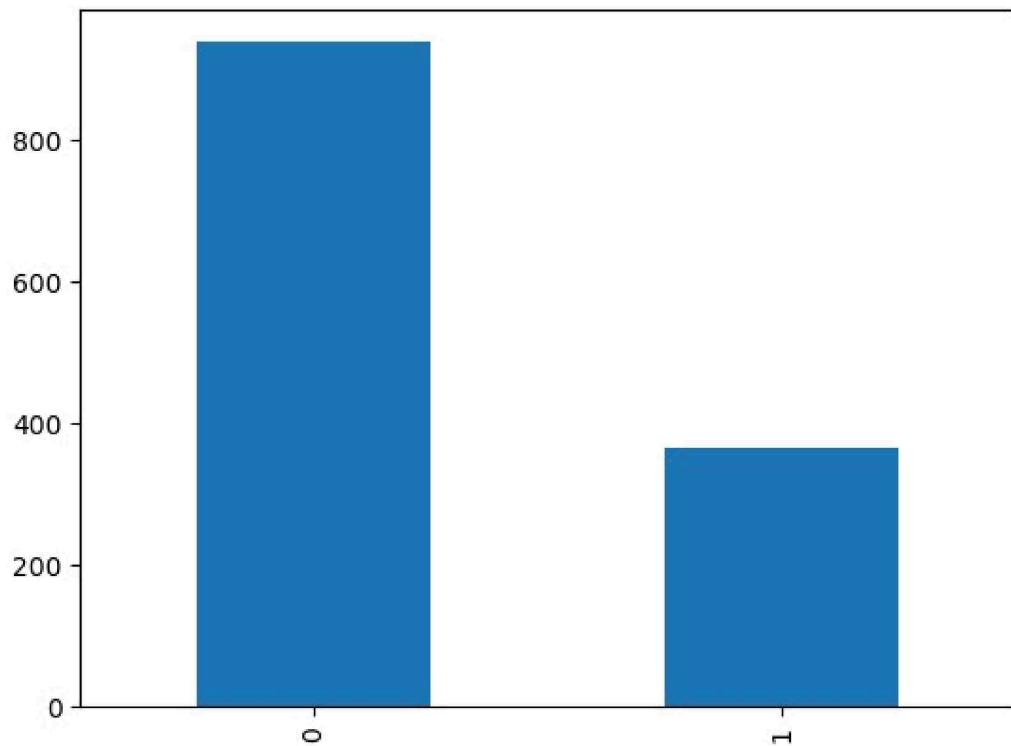
           Gpu  OpSys  Weight      Price
Touchscreen  Ips
0  Intel Iris Plus Graphics 640  macOS    1.37  71378.6832
0      1
1  Intel HD Graphics 6000  macOS    1.34  47895.5232
0      0

```

```
2      Intel HD Graphics 620  No OS     1.86  30636.0000
0      0
3      AMD Radeon Pro 455  macOS     1.83  135195.3360
0      1
4  Intel Iris Plus Graphics 650  macOS     1.37   96095.8080
0      1

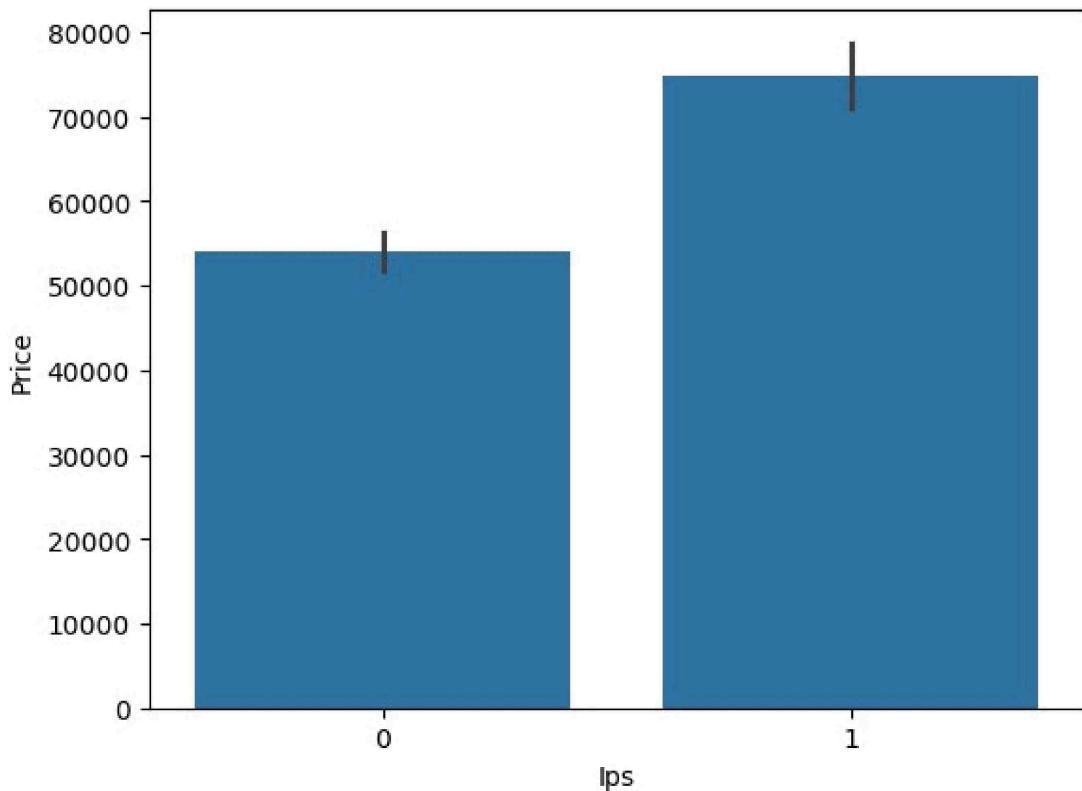
df['Ips'].value_counts().plot(kind='bar')

<AxesSubplot:>
```



```
sns.barplot(x=df['Ips'],y=df['Price'])

<AxesSubplot:xlabel='Ips', ylabel='Price'>
```



```

new = df['ScreenResolution'].str.split('x',n=1,expand=True)

df['X_res'] = new[0]
df['Y_res'] = new[1]

df.sample(5)

      Company    TypeName   Inches           ScreenResolution \
536      Dell     Notebook    15.6          Full HD 1920x1080
1007      HP      Ultrabook    14.0          Full HD 1920x1080
1244      HP     Notebook    14.0          Full HD 1920x1080
469      Lenovo    Ultrabook    14.0  IPS Panel Quad HD+ 2560x1440
363      HP     Notebook    15.6          Full HD 1920x1080

                                         Cpu   Ram       Memory
Gpu \
536      Intel Core i5 8250U 1.6GHz     8  256GB SSD        AMD Radeon
530
1007      Intel Core i7 6600U 2.6GHz     8  256GB SSD  Intel HD Graphics
520
1244      Intel Core i5 6200U 2.3GHz     4  256GB SSD  Intel HD Graphics
520
469      Intel Core i7 6600U 2.6GHz    12  256GB SSD  Intel HD Graphics
520
363      Intel Core i5 7200U 2.5GHz     8      1TB HDD  Intel HD Graphics

```

620

	OpSys	Weight	Price	Touchscreen	Ips	\
536	Windows 10	2.20	42486.0048	0	0	
1007	Windows 7	1.43	77202.7200	0	0	
1244	Windows 7	1.54	58607.4672	0	0	
469	Windows 10	1.40	98994.2400	0	1	
363	Windows 10	1.86	34045.9200	0	0	

	X_res	Y_res
536	Full HD	1920 1080
1007	Full HD	1920 1080
1244	Full HD	1920 1080
469	IPS Panel Quad HD+	2560 1440
363	Full HD	1920 1080

```
df['X_res'] = df['X_res'].str.replace(',', '').str.findall(r'(\d+\.\?|\d+)').apply(lambda x:x[0])
```

df.head()

	Company	TypeName	Inches	ScreenResolution	\
0	Apple	Ultrabook	13.3	IPS Panel Retina Display	2560x1600
1	Apple	Ultrabook	13.3		1440x900
2	HP	Notebook	15.6		Full HD 1920x1080
3	Apple	Ultrabook	15.4	IPS Panel Retina Display	2880x1800
4	Apple	Ultrabook	13.3	IPS Panel Retina Display	2560x1600

	Cpu	Ram	Memory	\
0	Intel Core i5 2.3GHz	8	128GB SSD	
1	Intel Core i5 1.8GHz	8	128GB Flash Storage	
2	Intel Core i5 7200U 2.5GHz	8	256GB SSD	
3	Intel Core i7 2.7GHz	16	512GB SSD	
4	Intel Core i5 3.1GHz	8	256GB SSD	

	Gpu	OpSys	Weight	Price
Touchscreen	Ips	\		
0	Intel Iris Plus Graphics 640	macOS	1.37	71378.6832
0	1			
1	Intel HD Graphics 6000	macOS	1.34	47895.5232
0	0			
2	Intel HD Graphics 620	No OS	1.86	30636.0000
0	0			
3	AMD Radeon Pro 455	macOS	1.83	135195.3360
0	1			
4	Intel Iris Plus Graphics 650	macOS	1.37	96095.8080
0	1			

	X_res	Y_res
0	2560	1600

```

1 1440 900
2 1920 1080
3 2880 1800
4 2560 1600

df['X_res'] = df['X_res'].astype('int')
df['Y_res'] = df['Y_res'].astype('int')

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1303 entries, 0 to 1302
Data columns (total 15 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Company          1303 non-null    object  
 1   TypeName         1303 non-null    object  
 2   Inches           1303 non-null    float64 
 3   ScreenResolution 1303 non-null    object  
 4   Cpu              1303 non-null    object  
 5   Ram              1303 non-null    int32   
 6   Memory           1303 non-null    object  
 7   Gpu              1303 non-null    object  
 8   OpSys            1303 non-null    object  
 9   Weight            1303 non-null    float32  
 10  Price             1303 non-null    float64 
 11  Touchscreen       1303 non-null    int64   
 12  Ips              1303 non-null    int64   
 13  X_res            1303 non-null    int32   
 14  Y_res            1303 non-null    int32   
dtypes: float32(1), float64(2), int32(3), int64(2), object(7)
memory usage: 96.8+ KB

df.corr()['Price']

<ipython-input-41-9447c1bc3d29>:1: FutureWarning: The default value of
numeric_only in DataFrame.corr is deprecated. In a future version, it
will default to False. Select only valid columns or specify the value
of numeric_only to silence this warning.
df.corr()['Price']

Inches          0.068197
Ram             0.743007
Weight          0.210370
Price           1.000000
Touchscreen     0.191226
Ips             0.252208
X_res           0.556529
Y_res           0.552809
Name: Price, dtype: float64

```

```

df['ppi'] = (((df['X_res']**2) +
(df['Y_res']**2))**0.5/df['Inches']).astype('float')

df.corr()['Price']

<ipython-input-43-9447c1bc3d29>:1: FutureWarning: The default value of
numeric_only in DataFrame.corr is deprecated. In a future version, it
will default to False. Select only valid columns or specify the value
of numeric_only to silence this warning.
df.corr()['Price']

Inches      0.068197
Ram         0.743007
Weight       0.210370
Price        1.000000
Touchscreen   0.191226
Ips          0.252208
X_res        0.556529
Y_res        0.552809
ppi          0.473487
Name: Price, dtype: float64

df.drop(columns=['ScreenResolution'], inplace=True)

df.head()

    Company    TypeName  Inches           Cpu  Ram \
0     Apple    Ultrabook    13.3  Intel Core i5 2.3GHz    8
1     Apple    Ultrabook    13.3  Intel Core i5 1.8GHz    8
2      HP    Notebook    15.6  Intel Core i5 7200U 2.5GHz    8
3     Apple    Ultrabook    15.4  Intel Core i7 2.7GHz   16
4     Apple    Ultrabook    13.3  Intel Core i5 3.1GHz    8

    Memory           Gpu  OpSys  Weight \
0  128GB SSD  Intel Iris Plus Graphics 640  macOS   1.37
1  128GB Flash Storage  Intel HD Graphics 6000  macOS   1.34
2  256GB SSD    Intel HD Graphics 620  No OS   1.86
3  512GB SSD    AMD Radeon Pro 455  macOS   1.83
4  256GB SSD  Intel Iris Plus Graphics 650  macOS   1.37

    Price  Touchscreen  Ips  X_res  Y_res      ppi
0  71378.6832          0    1   2560   1600  226.983005
1  47895.5232          0    0   1440    900  127.677940
2  30636.0000          0    0   1920   1080  141.211998
3  135195.3360          0    1   2880   1800  220.534624
4  96095.8080          0    1   2560   1600  226.983005

df.drop(columns=['Inches', 'X_res', 'Y_res'], inplace=True)

df.head()

```

```

Company    TypeName          Cpu   Ram
Memory \
0 Apple Ultrabook        Intel Core i5 2.3GHz  8      128GB
SSD
1 Apple Ultrabook        Intel Core i5 1.8GHz  8      128GB Flash
Storage
2 HP   Notebook   Intel Core i5 7200U 2.5GHz  8      256GB
SSD
3 Apple Ultrabook        Intel Core i7 2.7GHz  16     512GB
SSD
4 Apple Ultrabook        Intel Core i5 3.1GHz  8      256GB
SSD

Gpu   OpSys  Weight      Price
Touchscreen Ips \
0 Intel Iris Plus Graphics 640  macOS    1.37    71378.6832
0   1
1       Intel HD Graphics 6000  macOS    1.34    47895.5232
0   0
2       Intel HD Graphics 620   No OS   1.86    30636.0000
0   0
3       AMD Radeon Pro 455   macOS    1.83    135195.3360
0   1
4 Intel Iris Plus Graphics 650  macOS    1.37    96095.8080
0   1

ppi
0 226.983005
1 127.677940
2 141.211998
3 220.534624
4 226.983005

df['Cpu'].value_counts()

Intel Core i5 7200U 2.5GHz      190
Intel Core i7 7700HQ 2.8GHz     146
Intel Core i7 7500U 2.7GHz     134
Intel Core i7 8550U 1.8GHz      73
Intel Core i5 8250U 1.6GHz      72
...
Intel Core M M3-6Y30 0.9GHz      1
AMD A9-Series 9420 2.9GHz       1
Intel Core i3 6006U 2.2GHz       1
AMD A6-Series 7310 2GHz         1
Intel Xeon E3-1535M v6 3.1GHz    1
Name: Cpu, Length: 118, dtype: int64

df['Cpu Name'] = df['Cpu'].apply(lambda x:" ".join(x.split()[0:3]))

```

```

df.head()

   Company    TypeName          Cpu   Ram
Memory \
0   Apple     Ultrabook  Intel Core i5 2.3GHz    8      128GB
SSD
1   Apple     Ultrabook  Intel Core i5 1.8GHz    8  128GB Flash
Storage
2     HP     Notebook  Intel Core i5 7200U 2.5GHz    8      256GB
SSD
3   Apple     Ultrabook  Intel Core i7 2.7GHz   16      512GB
SSD
4   Apple     Ultrabook  Intel Core i5 3.1GHz    8      256GB
SSD

   Gpu  OpSys  Weight        Price
Touchscreen Ips \
0  Intel Iris Plus Graphics 640  macOS    1.37  71378.6832
0     1
1  Intel HD Graphics 6000  macOS    1.34  47895.5232
0     0
2  Intel HD Graphics 620  No OS    1.86  30636.0000
0     0
3  AMD Radeon Pro 455  macOS    1.83  135195.3360
0     1
4  Intel Iris Plus Graphics 650  macOS    1.37  96095.8080
0     1

   ppi      Cpu Name
0  226.983005  Intel Core i5
1  127.677940  Intel Core i5
2  141.211998  Intel Core i5
3  220.534624  Intel Core i7
4  226.983005  Intel Core i5

def fetch_processor(text):
    if text == 'Intel Core i7' or text == 'Intel Core i5' or text ==
'IIntel Core i3':
        return text
    else:
        if text.split()[0] == 'Intel':
            return 'Other Intel Processor'
        else:
            return 'AMD Processor'

df['Cpu brand'] = df['Cpu Name'].apply(fetch_processor)
df.head()

   Company    TypeName          Cpu   Ram
Memory \

```

```

0 Apple Ultrabook           Intel Core i5 2.3GHz    8          128GB
SSD
1 Apple Ultrabook           Intel Core i5 1.8GHz    8  128GB Flash
Storage
2 HP Notebook   Intel Core i5 7200U 2.5GHz    8          256GB
SSD
3 Apple Ultrabook           Intel Core i7 2.7GHz    16         512GB
SSD
4 Apple Ultrabook           Intel Core i5 3.1GHz    8          256GB
SSD

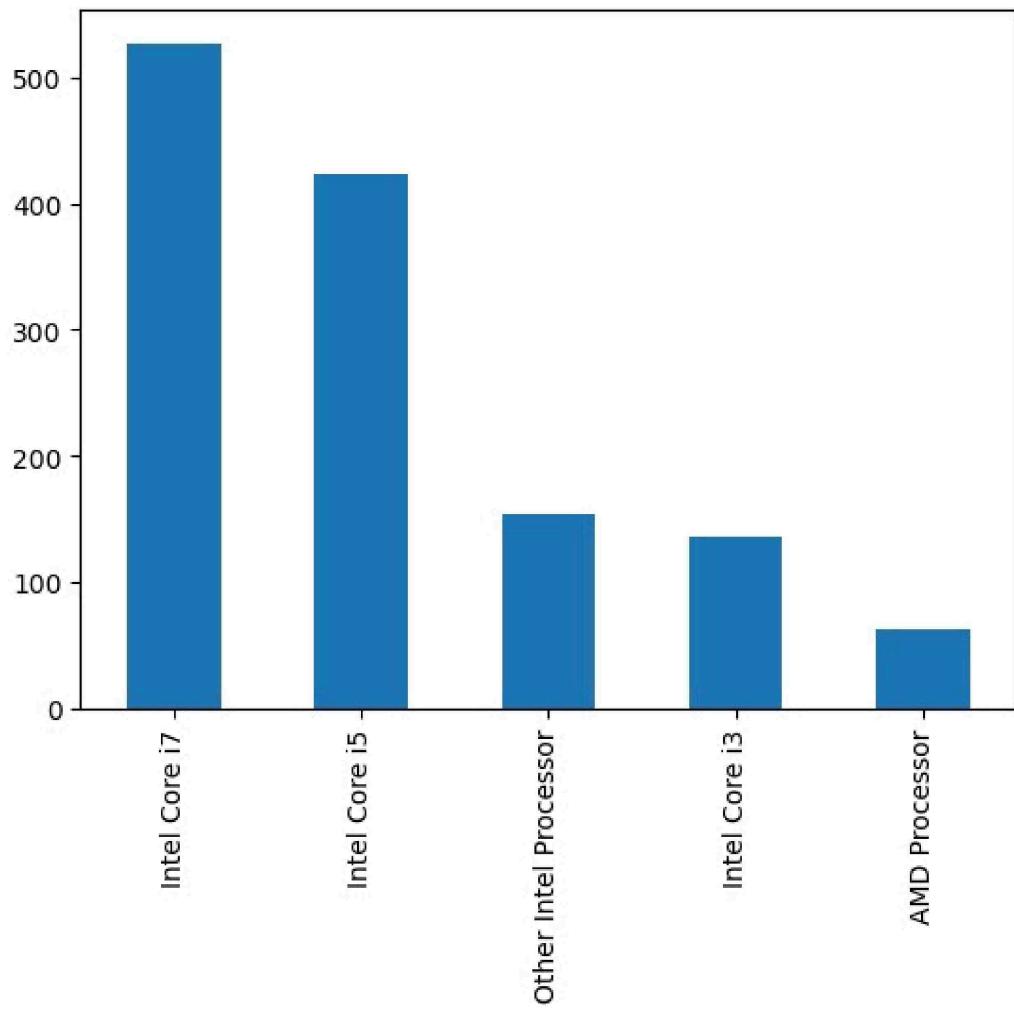
          Gpu  OpSys  Weight      Price
Touchscreen Ips \
0 Intel Iris Plus Graphics 640  macOS     1.37  71378.6832
0 1
1 Intel HD Graphics 6000  macOS     1.34  47895.5232
0 0
2 Intel HD Graphics 620   No OS     1.86  30636.0000
0 0
3 AMD Radeon Pro 455    macOS     1.83  135195.3360
0 1
4 Intel Iris Plus Graphics 650  macOS     1.37  96095.8080
0 1

      ppi      Cpu Name      Cpu brand
0 226.983005 Intel Core i5  Intel Core i5
1 127.677940 Intel Core i5  Intel Core i5
2 141.211998 Intel Core i5  Intel Core i5
3 220.534624 Intel Core i7  Intel Core i7
4 226.983005 Intel Core i5  Intel Core i5

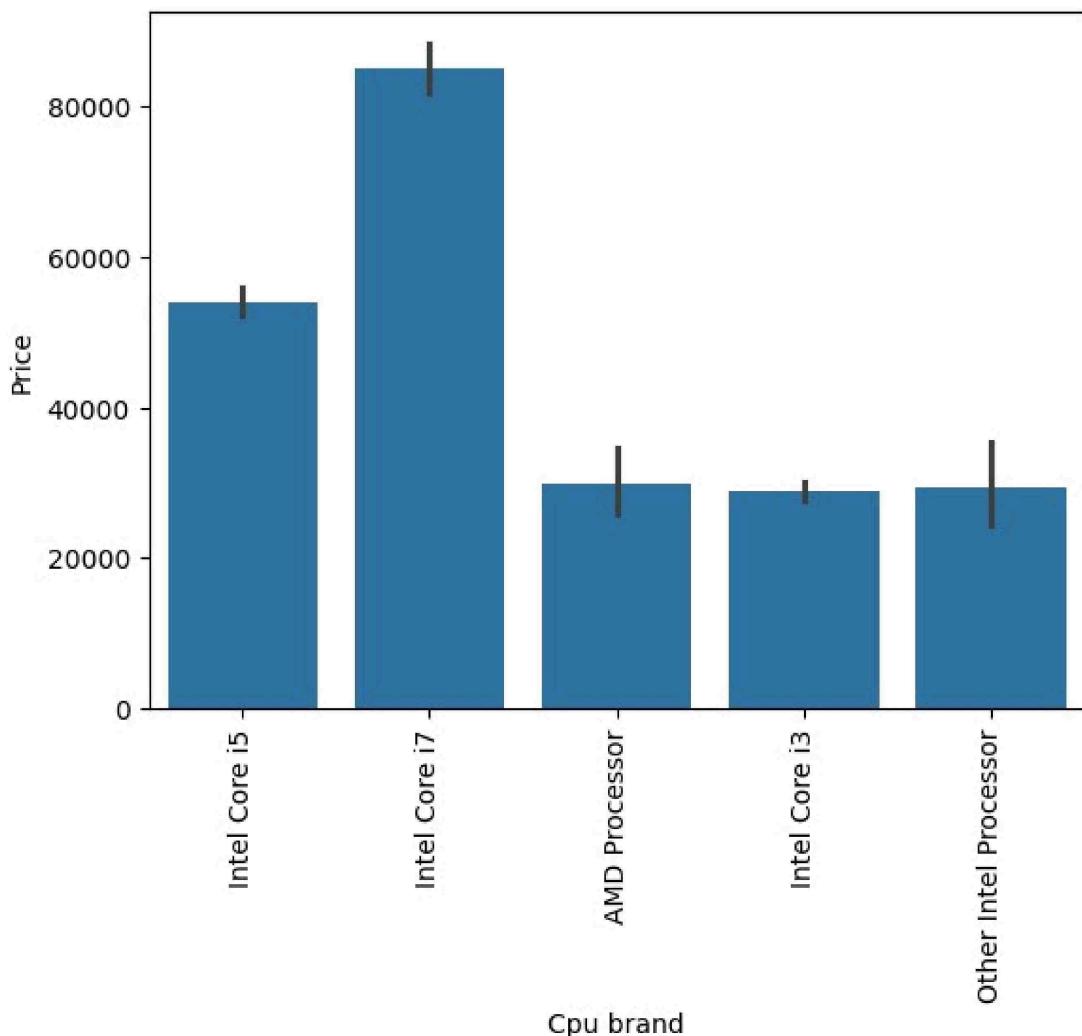
df['Cpu brand'].value_counts().plot(kind='bar')

<AxesSubplot:>

```



```
sns.barplot(x=df['Cpu brand'],y=df['Price'])
plt.xticks(rotation='vertical')
plt.show()
```



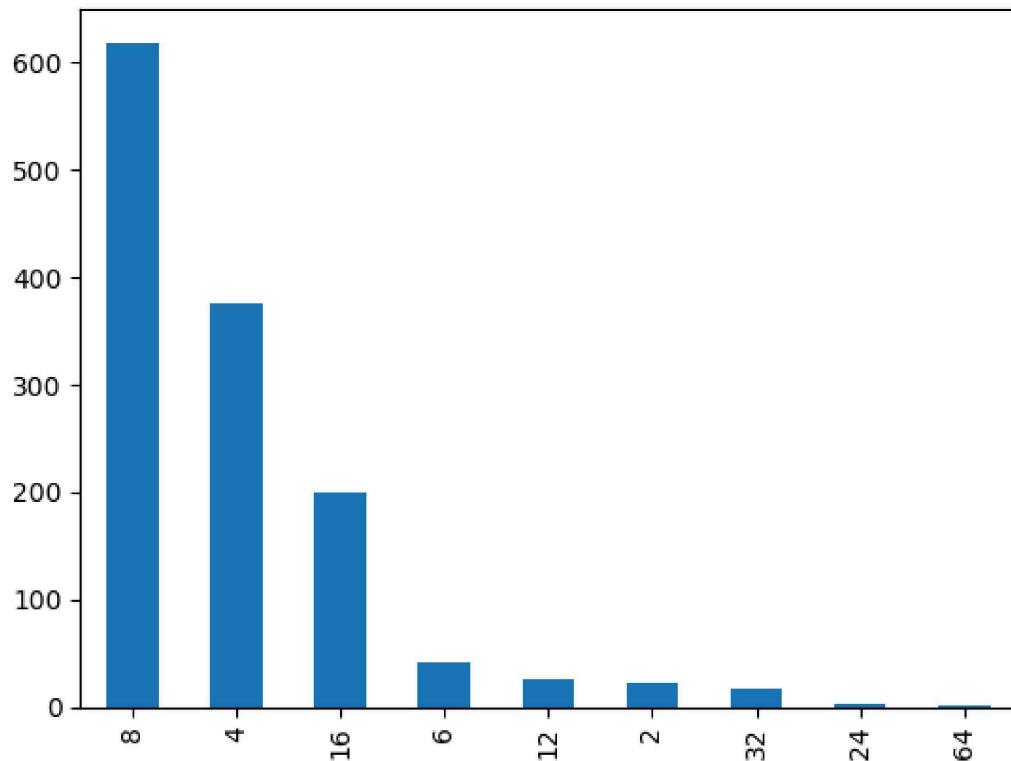
```
df.drop(columns=[ 'Cpu' , 'Cpu Name' ] ,inplace=True)
df.head()

   Company    TypeName  Ram          Memory
Gpu \
0   Apple     Ultrabook   8      128GB SSD  Intel Iris Plus
Graphics 640
1   Apple     Ultrabook   8  128GB Flash Storage  Intel HD
Graphics 6000
2     HP     Notebook   8      256GB SSD  Intel HD
Graphics 620
3   Apple     Ultrabook  16      512GB SSD  AMD Radeon
Pro 455
4   Apple     Ultrabook   8      256GB SSD  Intel Iris Plus
Graphics 650
```

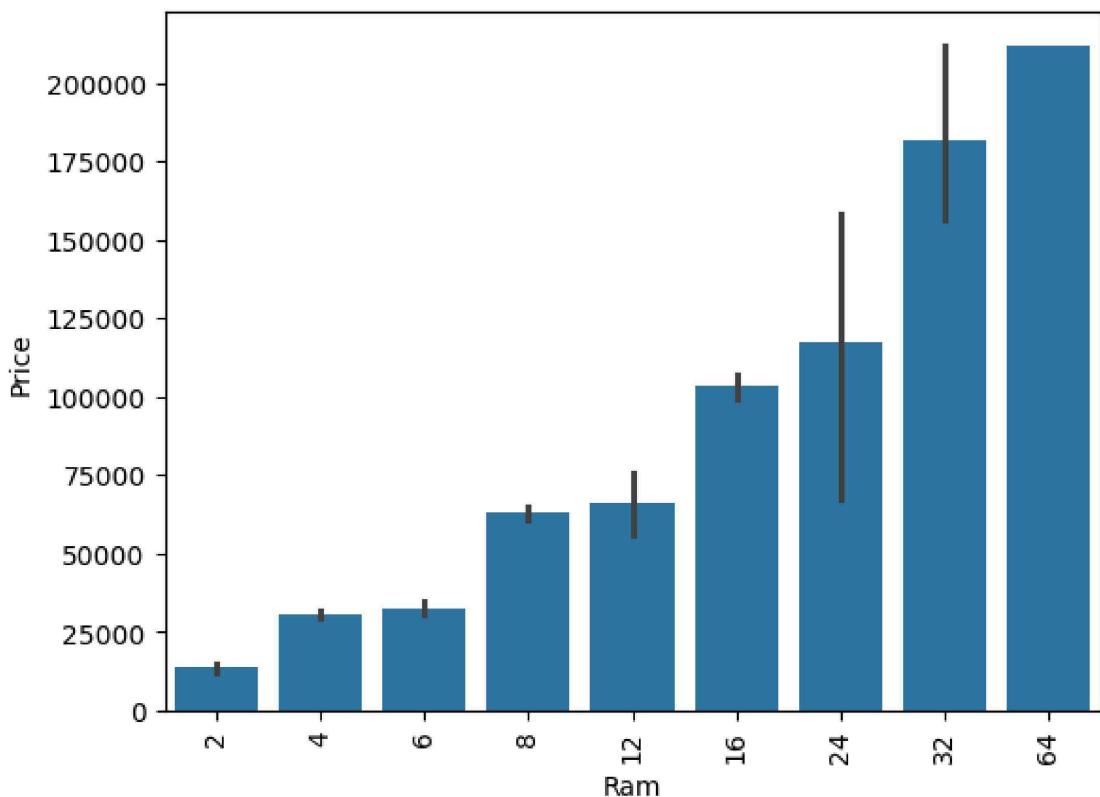
```
OpSys Weight Price Touchscreen Ips ppi Cpu  
brand  
0 macOS 1.37 71378.6832 0 1 226.983005 Intel  
Core i5  
1 macOS 1.34 47895.5232 0 0 127.677940 Intel  
Core i5  
2 No OS 1.86 30636.0000 0 0 141.211998 Intel  
Core i5  
3 macOS 1.83 135195.3360 0 1 220.534624 Intel  
Core i7  
4 macOS 1.37 96095.8080 0 1 226.983005 Intel  
Core i5
```

```
df['Ram'].value_counts().plot(kind='bar')
```

```
<AxesSubplot:>
```



```
sns.barplot(x=df['Ram'],y=df['Price'])  
plt.xticks(rotation='vertical')  
plt.show()
```



```
df[ 'Memory' ].value_counts()
```

256GB SSD	412
1TB HDD	223
500GB HDD	132
512GB SSD	118
128GB SSD + 1TB HDD	94
128GB SSD	76
256GB SSD + 1TB HDD	73
32GB Flash Storage	38
2TB HDD	16
64GB Flash Storage	15
512GB SSD + 1TB HDD	14
1TB SSD	14
256GB SSD + 2TB HDD	10
1.0TB Hybrid	9
256GB Flash Storage	8
16GB Flash Storage	7
32GB SSD	6
180GB SSD	5
128GB Flash Storage	4
512GB SSD + 2TB HDD	3
16GB SSD	3
512GB Flash Storage	2

```

1TB SSD + 1TB HDD          2
256GB SSD + 500GB HDD       2
128GB SSD + 2TB HDD         2
256GB SSD + 256GB SSD       2
512GB SSD + 256GB SSD       1
512GB SSD + 512GB SSD       1
64GB Flash Storage + 1TB HDD 1
1TB HDD + 1TB HDD           1
32GB HDD                     1
64GB SSD                     1
128GB HDD                     1
240GB SSD                     1
8GB SSD                      1
508GB Hybrid                  1
1.0TB HDD                     1
512GB SSD + 1.0TB Hybrid      1
256GB SSD + 1.0TB Hybrid      1
Name: Memory, dtype: int64

df['Memory'] = df['Memory'].astype(str).replace('.0', '', regex=True)
df["Memory"] = df["Memory"].str.replace('GB', '')
df["Memory"] = df["Memory"].str.replace('TB', '000')
new = df["Memory"].str.split("+", n = 1, expand = True)

df["first"] = new[0]
df["first"] = df["first"].str.strip()

df["second"] = new[1]

df["Layer1HDD"] = df["first"].apply(lambda x: 1 if "HDD" in x else 0)
df["Layer1SSD"] = df["first"].apply(lambda x: 1 if "SSD" in x else 0)
df["Layer1Hybrid"] = df["first"].apply(lambda x: 1 if "Hybrid" in x
else 0)
df["Layer1Flash_Storage"] = df["first"].apply(lambda x: 1 if "Flash
Storage" in x else 0)

df['first'] = df['first'].str.replace(r'\D', '')

df["second"].fillna("0", inplace = True)

df["Layer2HDD"] = df["second"].apply(lambda x: 1 if "HDD" in x else 0)
df["Layer2SSD"] = df["second"].apply(lambda x: 1 if "SSD" in x else 0)
df["Layer2Hybrid"] = df["second"].apply(lambda x: 1 if "Hybrid" in x
else 0)
df["Layer2Flash_Storage"] = df["second"].apply(lambda x: 1 if "Flash
Storage" in x else 0)

df['second'] = df['second'].str.replace(r'\D', '')

df["first"] = df["first"].astype(int)

```

```

df["second"] = df["second"].astype(int)

df["HDD"]=(df["first"]*df["Layer1HDD"]+df["second"]*df["Layer2HDD"])
df["SSD"]=(df["first"]*df["Layer1SSD"]+df["second"]*df["Layer2SSD"])
df["Hybrid"]=(df["first"]*df["Layer1Hybrid"]
+df["second"]*df["Layer2Hybrid"])
df["Flash_Storage"]=(df["first"]*df["Layer1Flash_Storage"]
+df["second"]*df["Layer2Flash_Storage"])

df.drop(columns=['first', 'second', 'Layer1HDD', 'Layer1SSD',
'Layer1Hybrid',
'Layer1Flash_Storage', 'Layer2HDD', 'Layer2SSD',
'Layer2Hybrid',
'Layer2Flash_Storage'],inplace=True)

```

```

<ipython-input-61-10829db803de>:16: FutureWarning: The default value
of regex will change from True to False in a future version.
    df['first'] = df['first'].str.replace(r'\D', '')
<ipython-input-61-10829db803de>:25: FutureWarning: The default value
of regex will change from True to False in a future version.
    df['second'] = df['second'].str.replace(r'\D', '')

```

```
df.sample(5)
```

Gpu	Company	Type	Name	Ram	Memory		
1155	HP	Notebook		4	256	SSD	AMD Radeon
520							
538	HP	Gaming		12	1000	HDD	Nvidia GeForce GTX
1060							
837	Toshiba	Notebook		4	500	HDD	Intel HD Graphics
620							
445	Dell	Notebook		8	256	SSD	Nvidia GeForce
930MX							
955	Dell	Gaming		16	512	SSD + 1000 HDD	Nvidia GeForce GTX
1070							

	OpSys	Weight	Price	Touchscreen	Ips	ppi	\
1155	Windows 10	1.91	25515.2592	0	0	141.211998	
538	Windows 10	2.62	95850.7200	0	0	127.335675	
837	Windows 10	2.00	35644.3200	0	0	100.454670	
445	Windows 10	1.64	62817.1200	0	0	157.350512	
955	Windows 10	4.36	168045.1200	0	1	254.671349	

	Cpu	brand	HDD	SSD	Hybrid	Flash_Storage
1155	AMD	Processor	0	256	0	0
538	Intel	Core i7	1000	0	0	0
837	Intel	Core i3	500	0	0	0
445	Intel	Core i5	0	256	0	0
955	Intel	Core i7	1000	512	0	0

```

df.drop(columns=[ 'Memory' ],inplace=True)

df.head()

  Company TypeName Ram Gpu OpSys Weight
0 Apple Ultrabook 8 Intel Iris Plus Graphics 640 macOS 1.37
1 Apple Ultrabook 8 Intel HD Graphics 6000 macOS 1.34
2 HP Notebook 8 Intel HD Graphics 620 No OS 1.86
3 Apple Ultrabook 16 AMD Radeon Pro 455 macOS 1.83
4 Apple Ultrabook 8 Intel Iris Plus Graphics 650 macOS 1.37

  Price Touchscreen Ips ppi Cpu brand HDD SSD
Hybrid \
0 71378.6832 0 1 226.983005 Intel Core i5 0 128
0
1 47895.5232 0 0 127.677940 Intel Core i5 0 0
0
2 30636.0000 0 0 141.211998 Intel Core i5 0 256
0
3 135195.3360 0 1 220.534624 Intel Core i7 0 512
0
4 96095.8080 0 1 226.983005 Intel Core i5 0 256
0

  Flash_Storage
0 0
1 128
2 0
3 0
4 0

df.corr()['Price']

<ipython-input-66-9447c1bc3d29>:1: FutureWarning: The default value of
numeric_only in DataFrame.corr is deprecated. In a future version, it
will default to False. Select only valid columns or specify the value
of numeric_only to silence this warning.
df.corr()['Price']

Ram 0.743007
Weight 0.210370
Price 1.000000
Touchscreen 0.191226
Ips 0.252208
ppi 0.473487

```

```

HDD           -0.096441
SSD            0.670799
Hybrid         0.007989
Flash_Storage -0.040511
Name: Price, dtype: float64

df.drop(columns=['Hybrid','Flash_Storage'],inplace=True)

df.head()

   Company    TypeName   Ram          Gpu      OpSys  Weight
0   Apple     Ultrabook    8  Intel Iris Plus Graphics 640  macOS    1.37
1   Apple     Ultrabook    8          Intel HD Graphics 6000  macOS    1.34
2     HP      Notebook    8          Intel HD Graphics 620  No OS    1.86
3   Apple     Ultrabook   16          AMD Radeon Pro 455  macOS    1.83
4   Apple     Ultrabook    8  Intel Iris Plus Graphics 650  macOS    1.37

      Price  Touchscreen   Ips       ppi      Cpu brand  HDD  SSD
0  71378.6832            0    1  226.983005  Intel Core i5    0  128
1  47895.5232            0    0  127.677940  Intel Core i5    0    0
2  30636.0000            0    0  141.211998  Intel Core i5    0  256
3  135195.3360           0    1  220.534624  Intel Core i7    0  512
4  96095.8080            0    1  226.983005  Intel Core i5    0  256

df['Gpu'].value_counts()

Intel HD Graphics 620        281
Intel HD Graphics 520        185
Intel UHD Graphics 620         68
Nvidia GeForce GTX 1050        66
Nvidia GeForce GTX 1060        48
...
AMD Radeon R5 520             1
AMD Radeon R7                 1
Intel HD Graphics 540             1
AMD Radeon 540                 1
ARM Mali T860 MP4                 1
Name: Gpu, Length: 110, dtype: int64

df['Gpu brand'] = df['Gpu'].apply(lambda x:x.split()[0])

```

```

df.head()

   Company TypeName Ram Gpu OpSys Weight
0 Apple Ultrabook 8 Intel Iris Plus Graphics 640 macOS 1.37
1 Apple Ultrabook 8 Intel HD Graphics 6000 macOS 1.34
2 HP Notebook 8 Intel HD Graphics 620 No OS 1.86
3 Apple Ultrabook 16 AMD Radeon Pro 455 macOS 1.83
4 Apple Ultrabook 8 Intel Iris Plus Graphics 650 macOS 1.37

   Price Touchscreen Ips ppi Cpu brand HDD SSD
0 71378.6832 0 1 226.983005 Intel Core i5 0 128
1 47895.5232 0 0 127.677940 Intel Core i5 0 0
2 30636.0000 0 0 141.211998 Intel Core i5 0 256
3 135195.3360 0 1 220.534624 Intel Core i7 0 512
4 96095.8080 0 1 226.983005 Intel Core i5 0 256

   Gpu brand
0 Intel
1 Intel
2 Intel
3 AMD
4 Intel

df['Gpu brand'].value_counts()

Intel    722
Nvidia   400
AMD     180
ARM      1
Name: Gpu brand, dtype: int64

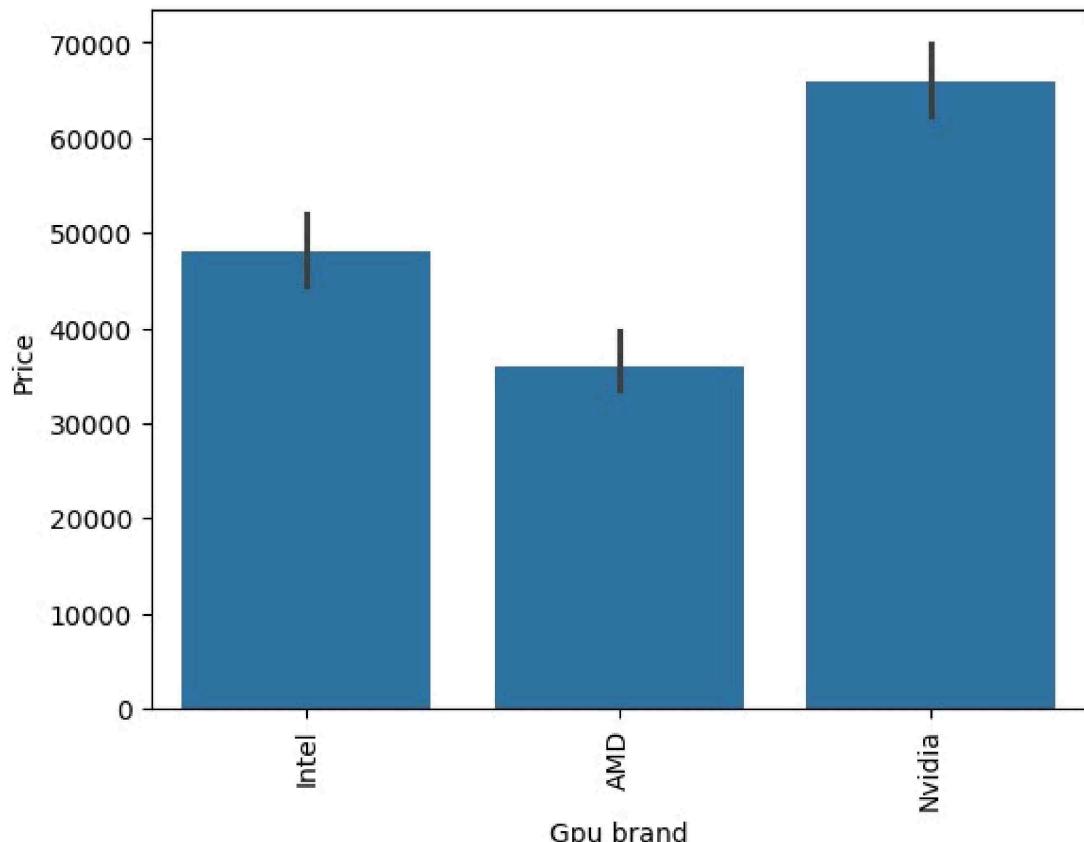
df = df[df['Gpu brand'] != 'ARM']

df['Gpu brand'].value_counts()

Intel    722
Nvidia   400
AMD     180
Name: Gpu brand, dtype: int64

```

```
sns.barplot(x=df['Gpu brand'],y=df['Price'],estimator=np.median)
plt.xticks(rotation='vertical')
plt.show()
```



```
df.drop(columns=[ 'Gpu'],inplace=True)

df.head()

Company TypeName Ram OpSys Weight Price Touchscreen
Ips \
0 Apple Ultrabook 8 macOS 1.37 71378.6832 0
1
1 Apple Ultrabook 8 macOS 1.34 47895.5232 0
0
2 HP Notebook 8 No OS 1.86 30636.0000 0
0
3 Apple Ultrabook 16 macOS 1.83 135195.3360 0
1
4 Apple Ultrabook 8 macOS 1.37 96095.8080 0
1
```

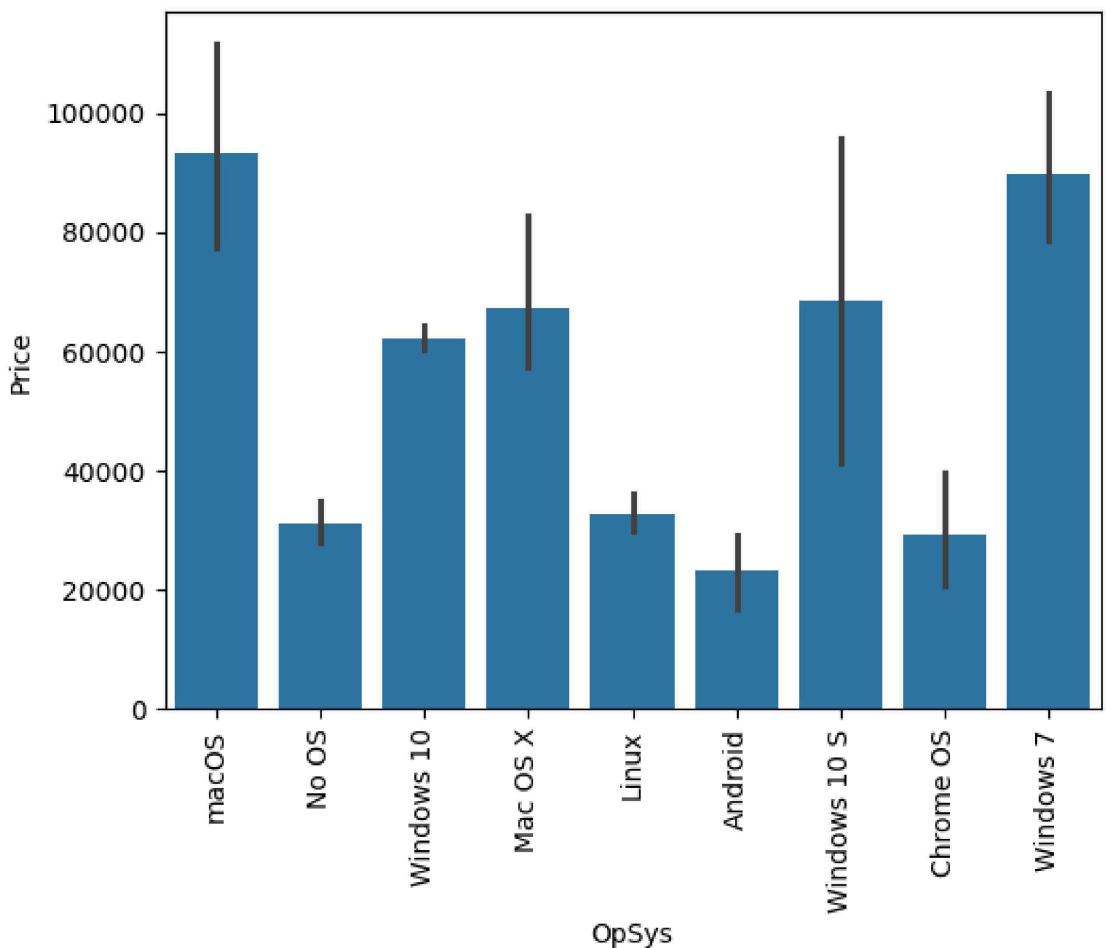
	ppi	Cpu brand	HDD	SSD	Gpu brand
--	-----	-----------	-----	-----	-----------

```
0 226.983005 Intel Core i5    0 128     Intel
1 127.677940 Intel Core i5    0 0       Intel
2 141.211998 Intel Core i5    0 256     Intel
3 220.534624 Intel Core i7    0 512     AMD
4 226.983005 Intel Core i5    0 256     Intel

df['OpSys'].value_counts()

Windows 10      1072
No OS           66
Linux            62
Windows 7        45
Chrome OS        26
macOS            13
Mac OS X          8
Windows 10 S       8
Android           2
Name: OpSys, dtype: int64

sns.barplot(x=df['OpSys'],y=df['Price'])
plt.xticks(rotation='vertical')
plt.show()
```



```
def cat_os(inp):
    if inp == 'Windows 10' or inp == 'Windows 7' or inp == 'Windows 10 S':
        return 'Windows'
    elif inp == 'macOS' or inp == 'Mac OS X':
        return 'Mac'
    else:
        return 'Others/No OS/Linux'

df['os'] = df['OpSys'].apply(cat_os)

df.head()

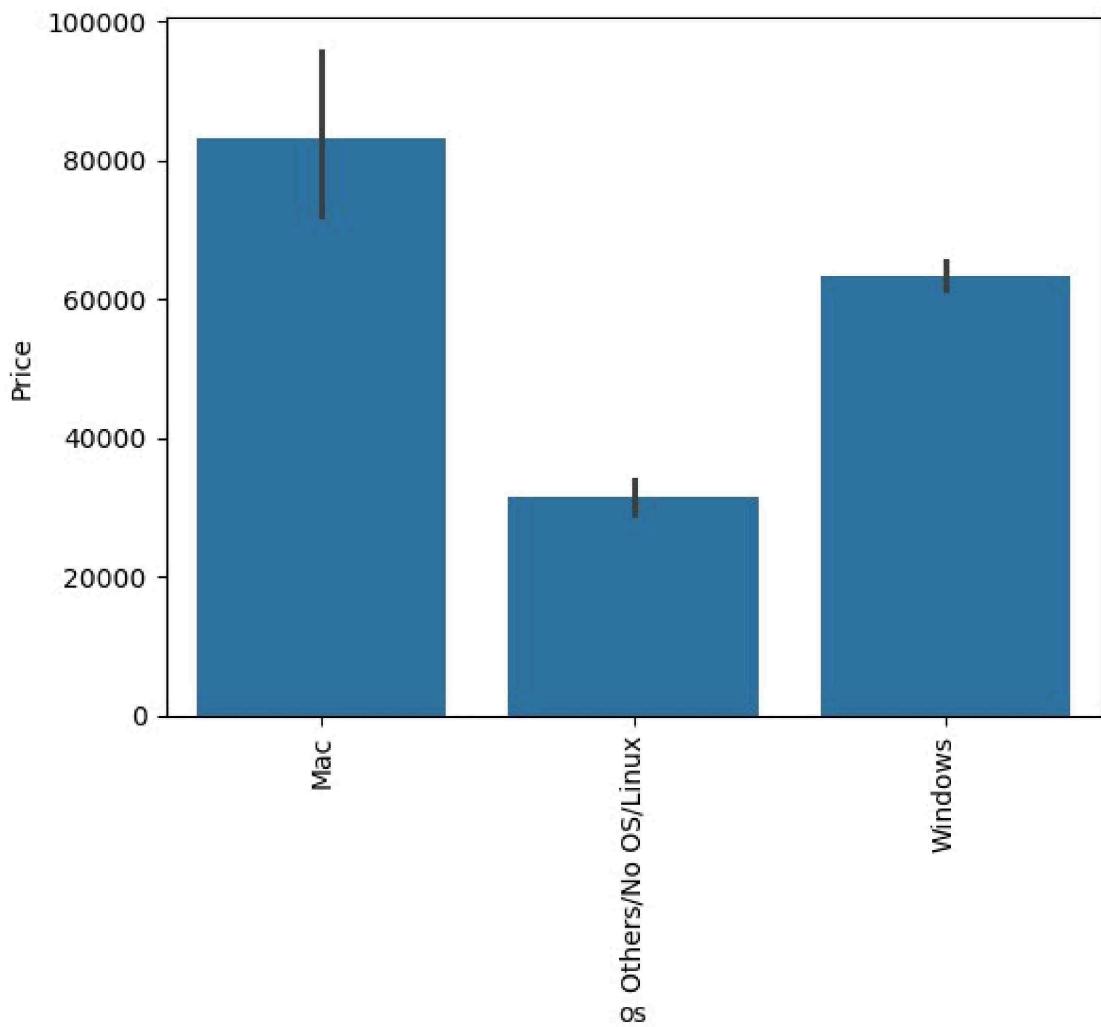
   Company   TypeName  Ram  OpSys  Weight      Price  Touchscreen
0     Apple  Ultrabook    8  macOS    1.37  71378.6832          0
1     Apple  Ultrabook    8  macOS    1.34  47895.5232          0
0      HP    Notebook    8  No OS    1.86  30636.0000          0
```

```
0
3   Apple Ultrabook    16  macOS      1.83  135195.3360          0
1
4   Apple Ultrabook     8  macOS      1.37   96095.8080          0
1

      ppi      Cpu brand  HDD  SSD Gpu brand           os
0  226.983005  Intel Core i5    0  128  Intel           Mac
1  127.677940  Intel Core i5    0    0  Intel           Mac
2  141.211998  Intel Core i5    0  256  Intel  Others/No OS/Linux
3  220.534624  Intel Core i7    0  512  AMD            Mac
4  226.983005  Intel Core i5    0  256  Intel           Mac

df.drop(columns=['OpSys'], inplace=True)

sns.barplot(x=df['os'], y=df['Price'])
plt.xticks(rotation='vertical')
plt.show()
```



```
sns.distplot(df['Weight'])

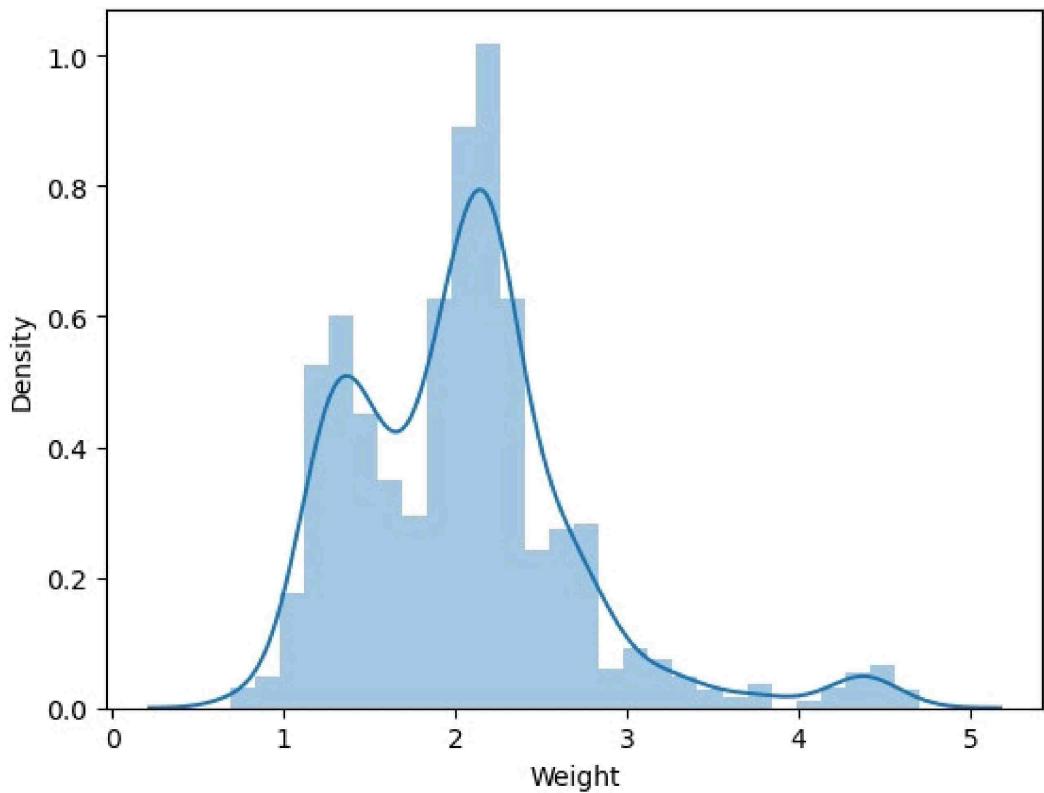
<ipython-input-85-05ee4c8848be>:1: UserWarning:
`distplot` is a deprecated function and will be removed in seaborn
v0.14.0.

Please adapt your code to use either `displot` (a figure-level
function with
similar flexibility) or `histplot` (an axes-level function for
histograms).

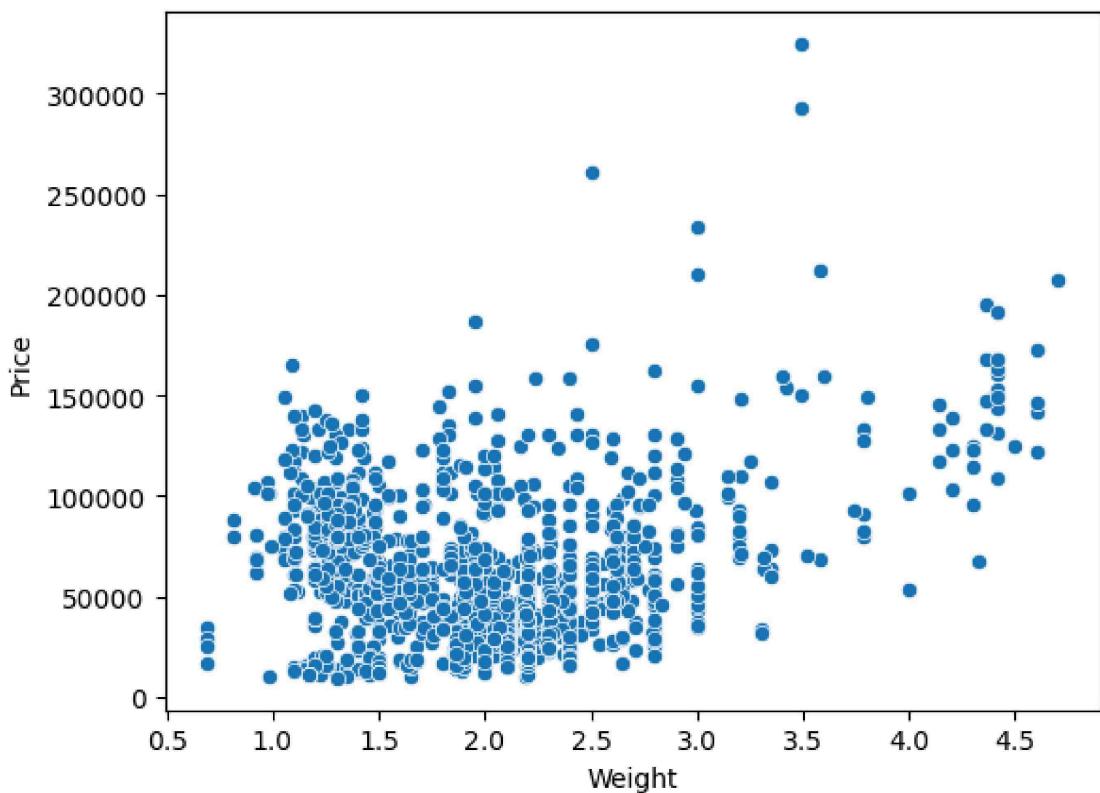
For a guide to updating your code to use the new functions, please see
https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751

sns.distplot(df['Weight'])

<AxesSubplot:xlabel='Weight', ylabel='Density'>
```



```
sns.scatterplot(x=df['Weight'],y=df['Price'])  
<AxesSubplot:xlabel='Weight', ylabel='Price'>
```



```
df.corr()['Price']

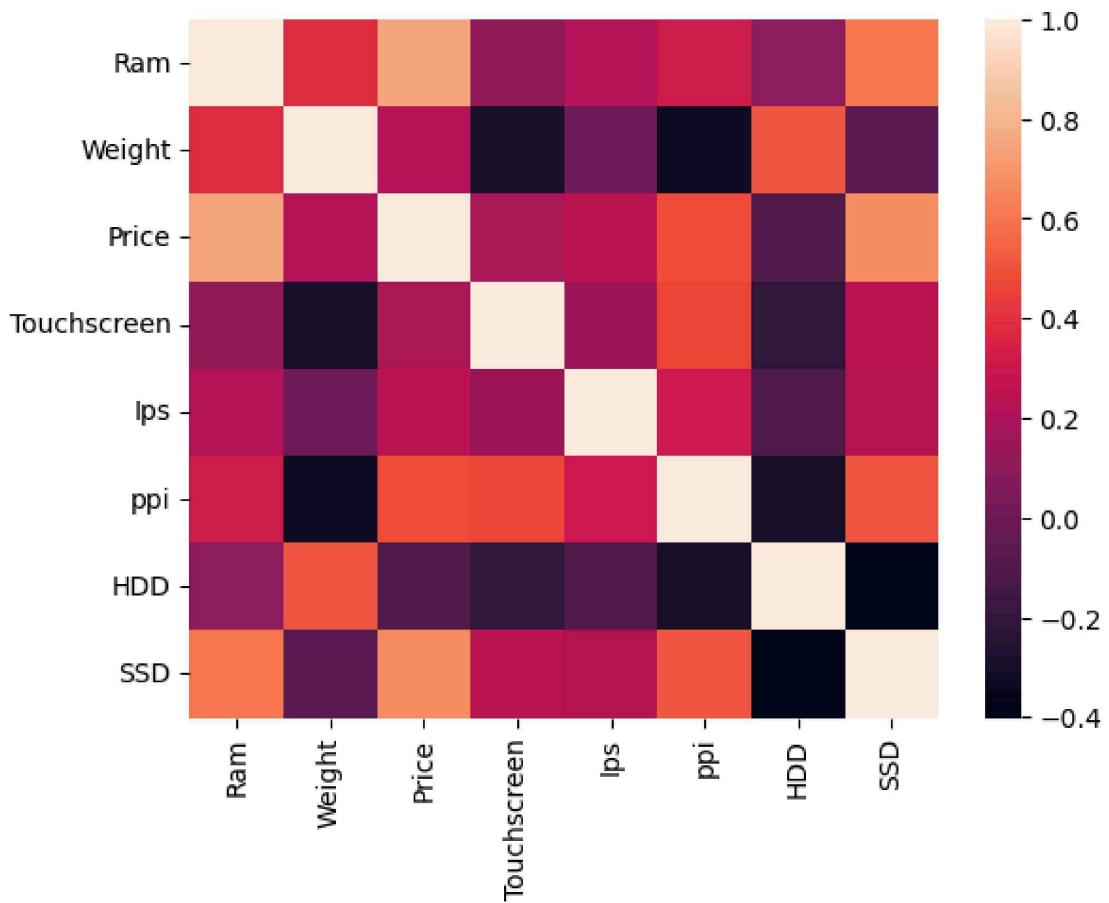
<ipython-input-87-9447c1bc3d29>:1: FutureWarning: The default value of
numeric_only in DataFrame.corr is deprecated. In a future version, it
will default to False. Select only valid columns or specify the value
of numeric_only to silence this warning.
    df.corr()['Price']

Ram          0.742905
Weight        0.209867
Price         1.000000
Touchscreen   0.192917
Ips           0.253320
ppi            0.475368
HDD           -0.096891
SSD            0.670660
Name: Price, dtype: float64

sns.heatmap(df.corr())

<ipython-input-88-aa4f4450a243>:1: FutureWarning: The default value of
numeric_only in DataFrame.corr is deprecated. In a future version, it
will default to False. Select only valid columns or specify the value
of numeric_only to silence this warning.
    sns.heatmap(df.corr())
```

```
<AxesSubplot:>
```



```
sns.distplot(np.log(df['Price']))
```

```
<ipython-input-89-c1a82a4801f0>:1: UserWarning:
```

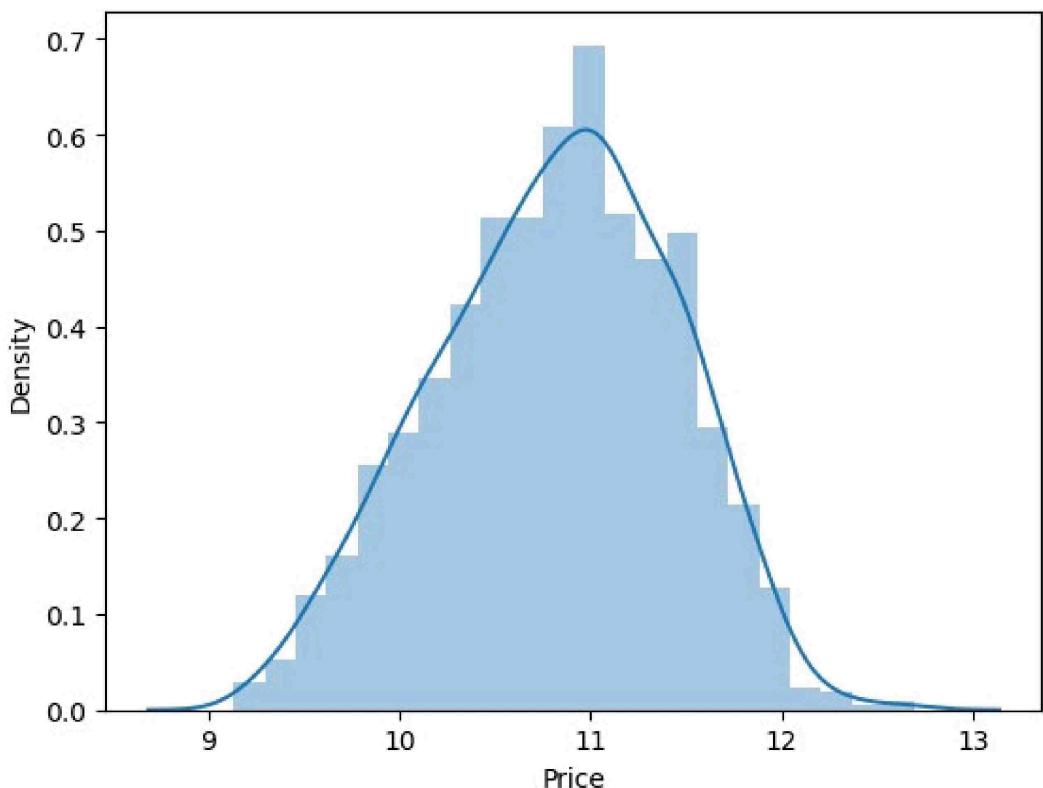
```
`distplot` is a deprecated function and will be removed in seaborn  
v0.14.0.
```

```
Please adapt your code to use either `displot` (a figure-level  
function with  
similar flexibility) or `histplot` (an axes-level function for  
histograms).
```

For a guide to updating your code to use the new functions, please see  
<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(np.log(df['Price']))
```

```
<AxesSubplot:xlabel='Price', ylabel='Density'>
```



```
X = df.drop(columns=['Price'])
y = np.log(df['Price'])

X
```

	Company	Type	ppi	TypeName	Ram	Weight	Touchscreen	Ips
0	Apple	Ultrabook	226.983005		8	1.37	0	1
1	Apple	Ultrabook	127.677940		8	1.34	0	0
2	HP	Notebook	141.211998		8	1.86	0	0
3	Apple	Ultrabook	220.534624		16	1.83	0	1
4	Apple	Ultrabook	226.983005		8	1.37	0	1
...	...	...	...	...	...	...	...	...
...	...	...	...	...	...	...	...	...
1298	Lenovo	2 in 1 Convertible	157.350512		4	1.80	1	1
1299	Lenovo	2 in 1 Convertible	276.053530		16	1.30	1	1
1300	Lenovo	Notebook			2	1.50	0	0

```

111.935204
1301      HP          Notebook   6    2.19          0    0
100.454670
1302      Asus         Notebook   4    2.20          0    0
100.454670

      Cpu brand  HDD  SSD Gpu brand           os
0      Intel Core i5  0  128  Intel           Mac
1      Intel Core i5  0   0  Intel           Mac
2      Intel Core i5  0  256  Intel  Others/No OS/Linux
3      Intel Core i7  0  512  AMD            Mac
4      Intel Core i5  0  256  Intel           Mac
...
1298     ...       ...  ...  ...           ...
1299     ...       ...  ...  ...           ...
1300  Other Intel Processor  0   0  Intel           Windows
1301     ...       ...  ...  ...           ...
1302  Other Intel Processor  500  0  Intel           Windows
[1302 rows x 12 columns]

y

0      11.175755
1      10.776777
2      10.329931
3      11.814476
4      11.473101
...
1298     10.433899
1299     11.288115
1300     9.409283
1301     10.614129
1302     9.886358
Name: Price, Length: 1302, dtype: float64

from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test =
train_test_split(X,y,test_size=0.15,random_state=2)

X_train

      Company        TypeName  Ram  Weight  Touchscreen  Ips
ppi \
183  Toshiba        Notebook   8    2.00          0    0
100.454670
1141  MSI           Gaming    8    2.40          0    0
141.211998
1049  Asus           Netbook   4    1.20          0    0
135.094211
1020  Dell  2 in 1 Convertible  4    2.08          1    1

```

```

141.211998
878     Dell          Notebook   4    2.18        0    0
141.211998
...
...
466     Acer          Notebook   4    2.20        0    0
100.454670
299     Asus          Ultrabook  16   1.63        0    0
141.211998
493     Acer          Notebook   8    2.20        0    0
100.454670
527     Lenovo         Notebook   8    2.20        0    0
100.454670
1193    Apple          Ultrabook  8    0.92        0    1
226.415547

           Cpu brand  HDD  SSD  Gpu brand      os
183       Intel Core i5  0    128  Intel  Windows
1141      Intel Core i7  1000 128  Nvidia Windows
1049  Other Intel Processor  0    0    Intel  Others/No OS/Linux
1020      Intel Core i3  1000 0    Intel  Windows
878       Intel Core i5  1000 128  Nvidia Windows
...
...
466       Intel Core i3  500  0    Nvidia Windows
299       Intel Core i7  0    512  Nvidia Windows
493       AMD Processor 1000 0    AMD   Windows
527       Intel Core i3  2000 0    Nvidia Others/No OS/Linux
1193  Other Intel Processor  0    0    Intel  Mac

```

[1106 rows x 12 columns]

```

from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import OneHotEncoder
from sklearn.metrics import r2_score, mean_absolute_error

from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.neighbors import KNeighborsRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import
RandomForestRegressor, GradientBoostingRegressor, AdaBoostRegressor, ExtraTreesRegressor
from sklearn.svm import SVR
from xgboost import XGBRegressor

```

## Linear regression

```

step1 = ColumnTransformer(transformers=[
    ('col_tnf', OneHotEncoder(sparse=False, drop='first'), [0, 1, 7, 10, 11]),
], remainder='passthrough')

```

```

step2 = LinearRegression()

pipe = Pipeline([
    ('step1',step1),
    ('step2',step2)
])
pipe.fit(X_train,y_train)

y_pred = pipe.predict(X_test)

print('R2 score',r2_score(y_test,y_pred))
print('MAE',mean_absolute_error(y_test,y_pred))

/lib/python3.11/site-packages/sklearn/preprocessing/_encoders.py:975:
FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2
and will be removed in 1.4. `sparse_output` is ignored unless you
leave `sparse` to its default value.

warnings.warn(
R2 score 0.8073277448418492
MAE 0.21017827976428904

```

## Random Forest

```

step1 = ColumnTransformer(transformers=[
    ('col_tnf',OneHotEncoder(sparse=False,drop='first'),[0,1,7,10,11])
],remainder='passthrough')

step2 = RandomForestRegressor(n_estimators=100,
                             random_state=3,
                             max_samples=0.5,
                             max_features=0.75,
                             max_depth=15)

pipe = Pipeline([
    ('step1',step1),
    ('step2',step2)
])
pipe.fit(X_train,y_train)

y_pred = pipe.predict(X_test)

print('R2 score',r2_score(y_test,y_pred))
print('MAE',mean_absolute_error(y_test,y_pred))

/lib/python3.11/site-packages/sklearn/preprocessing/_encoders.py:975:
FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2
and will be removed in 1.4. `sparse_output` is ignored unless you

```

```
leave `sparse` to its default value.  
warnings.warn(
```

```
R2 score 0.8873402378382488  
MAE 0.15860130110457718
```

## Exporting the Model

```
import pickle
```

```
pickle.dump(df,open('df.pkl','wb'))  
pickle.dump(pipe,open('pipe.pkl','wb'))
```

```
df
```

```
   Company    TypeName  Ram  Weight      Price  
Touchscreen  Ips  \\  
0     Apple        Ultrabook  8    1.37  71378.6832  
0     1           Ultrabook  8    1.34  47895.5232  
0     0           Ultrabook  8    1.86  30636.0000  
2     HP           Notebook  8    1.83  135195.3360  
0     0           Ultrabook  16   1.83  96095.8080  
4     Apple        Ultrabook  8    1.37  96095.8080  
0     1           Ultrabook  8    1.37  96095.8080  
...     ...       ...  ...  ...  
1298  Lenovo  2 in 1 Convertible  4    1.80  33992.6400  
1     1           Ultrabook  16   1.30  79866.7200  
1     1           Ultrabook  2    1.50  12201.1200  
1300  Lenovo        Notebook  6    2.19  40705.9200  
0     0           Ultrabook  4    2.20  19660.3200  
1302  Asus          Notebook  4    2.20  19660.3200  
0     0           Ultrabook  4    2.20  19660.3200  
  
          ppi          Cpu brand  HDD  SSD Gpu brand  \\  
0     226.983005      Intel Core i5  0    128  Intel  
1     127.677940      Intel Core i5  0     0  Intel  
2     141.211998      Intel Core i5  0    256  Intel  
3     220.534624      Intel Core i7  0    512  AMD  
4     226.983005      Intel Core i5  0    256  Intel  
...     ...       ...  ...  ...  
1298  157.350512      Intel Core i7  0    128  Intel  
1299  276.053530      Intel Core i7  0    512  Intel  
1300  111.935204  Other Intel Processor  0     0  Intel
```

```

1301 100.454670           Intel Core i7 1000 0          AMD
1302 100.454670 Other Intel Processor 500 0          Intel

```

```

          os
0          Mac
1          Mac
2    Others/No OS/Linux
3          Mac
4          Mac
...
1298      Windows
1299      Windows
1300      Windows
1301      Windows
1302      Windows

```

[1302 rows x 13 columns]

X\_train

	Company	TypeName	Ram	Weight	Touchscreen	Ips
ppi \						
183 Toshiba		Notebook	8	2.00	0	0
100.454670						
1141 MSI		Gaming	8	2.40	0	0
141.211998						
1049 Asus		Netbook	4	1.20	0	0
135.094211						
1020 Dell	2 in 1 Convertible		4	2.08	1	1
141.211998						
878 Dell		Notebook	4	2.18	0	0
141.211998						
...	...	...	...	...	...	...
...						
466 Acer		Notebook	4	2.20	0	0
100.454670						
299 Asus		Ultrabook	16	1.63	0	0
141.211998						
493 Acer		Notebook	8	2.20	0	0
100.454670						
527 Lenovo		Notebook	8	2.20	0	0
100.454670						
1193 Apple		Ultrabook	8	0.92	0	1
226.415547						

	Cpu brand	HDD	SSD	Gpu brand	os
183	Intel Core i5	0	128	Intel	Windows
1141	Intel Core i7	1000	128	Nvidia	Windows
1049	Other Intel Processor	0	0	Intel	Others/No OS/Linux
1020	Intel Core i3	1000	0	Intel	Windows

878	Intel Core i5	1000	128	Nvidia		Windows
...	...	...	...	...	...	...
466	Intel Core i3	500	0	Nvidia		Windows
299	Intel Core i7	0	512	Nvidia		Windows
493	AMD Processor	1000	0	AMD		Windows
527	Intel Core i3	2000	0	Nvidia	Others/No OS/Linux	
1193	Other Intel Processor	0	0	Intel		Mac

[1106 rows x 12 columns]

## Future Scope of Improvement

The future of your laptop price predictor model can be expanded in various ways to improve its accuracy, functionality, and user experience. Here are some potential areas for exploration:

### **Data Acquisition and Preprocessing:**

- Real-time Price Data: Integrate real-time pricing information from online retailers using web scraping techniques or APIs. This can help capture dynamic price fluctuations and provide more up-to-date predictions.
- Image Recognition: Explore using image recognition models to extract information from laptop product images, such as brand logos, design elements, or physical features that might influence price.

### **Model Development and Training:**

- Deep Learning Architectures: Implement deep learning models like Convolutional Neural Networks (CNNs) to learn complex patterns from laptop specification data and potentially achieve higher prediction accuracy.
- Transfer Learning: Leverage pre-trained deep learning models on image datasets (e.g., for image recognition of laptop features) and fine-tune them for your specific laptop price prediction task.
- Multi-output Models: Develop models that predict not just the overall price, but also component-wise prices (CPU, RAM, storage etc.). This can provide more granular insights into the price breakdown of a laptop.

### **Model Deployment and User Interface:**

- **Web Application:** Develop a user-friendly web application where users can easily input laptop specifications and receive predicted prices. This can make the model more accessible and user-friendly.
- **API Integration:** Create an API (Application Programming Interface) that allows other applications to integrate your price prediction functionality. This can expand the reach of your model to other platforms or services.
- **Price Alerts:** Implement a feature that allows users to set price alerts for specific laptops based on their predicted price range. This can be helpful for deal hunters who want to be notified when a desired laptop reaches their budgets

## Certificate

This is to certify that Mr. Gaurav Kumar of Asansol Engineering College, roll number: 10871022016, has successfully completed a project on *Prediction of Laptop Prices using Machine Learning with Python* under the guidance of Prof. Arnab Chakraborty.

-----  
Prof. Arnab Chakraborty  
Globsyn Finishing School

## Certificate

This is to certify that Mr. Subham Shaw of Asansol Engineering College, roll number: 10871022004, has successfully completed a project on *Prediction of Laptop Prices using Machine Learning with Python* under the guidance of Prof. Arnab Chakraborty.

-----  
Prof. Arnab Chakraborty  
Globsyn Finishing School

## Certificate

This is to certify that Mr. Nishant Yadav of Asansol Engineering College, roll number: 10871022029, has successfully completed a project on *Prediction of Laptop Prices using Machine Learning with Python* under the guidance of Prof. Arnab Chakraborty.

-----  
Prof. Arnab Chakraborty  
Globsyn Finishing School

## Certificate

This is to certify that Ms. Puja Verma of Asansol Engineering College, roll number: 10871022011, has successfully completed a project on *Prediction of Laptop Prices using Machine Learning with Python* under the guidance of Prof. Arnab Chakraborty.

-----  
Prof. Arnab Chakraborty  
Globsyn Finishing School

## Certificate

This is to certify that Ms. Angelina Deepshikha Hans of Asansol Engineering College, roll number: 10871022009, has successfully completed a project on *Prediction of Laptop Prices using Machine Learning with Python* under the guidance of Prof. Arnab Chakraborty.

-----  
Prof. Arnab Chakraborty  
Globsyn Finishing School